
SDODR – CLOUD MODELS

Version 05.09.2018 – OIH Architecture Call

Susanne Braun



The Fraunhofer-Gesellschaft at a Glance

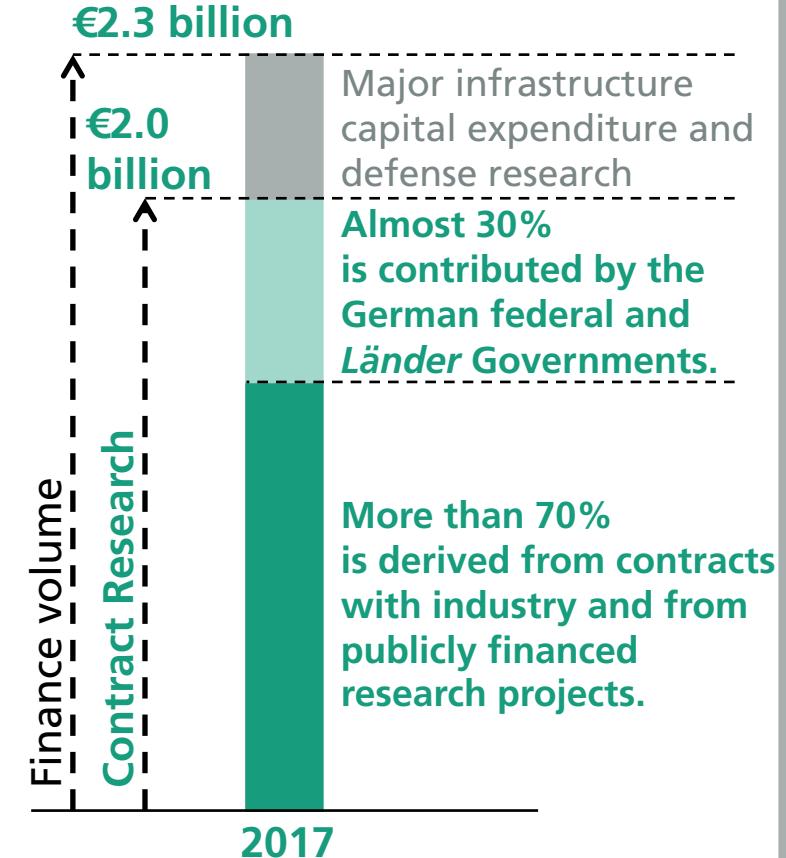
The Fraunhofer-Gesellschaft undertakes applied research of direct utility to private and public enterprise and of wide benefit to society.



25,527 staff



72 institutes and research units

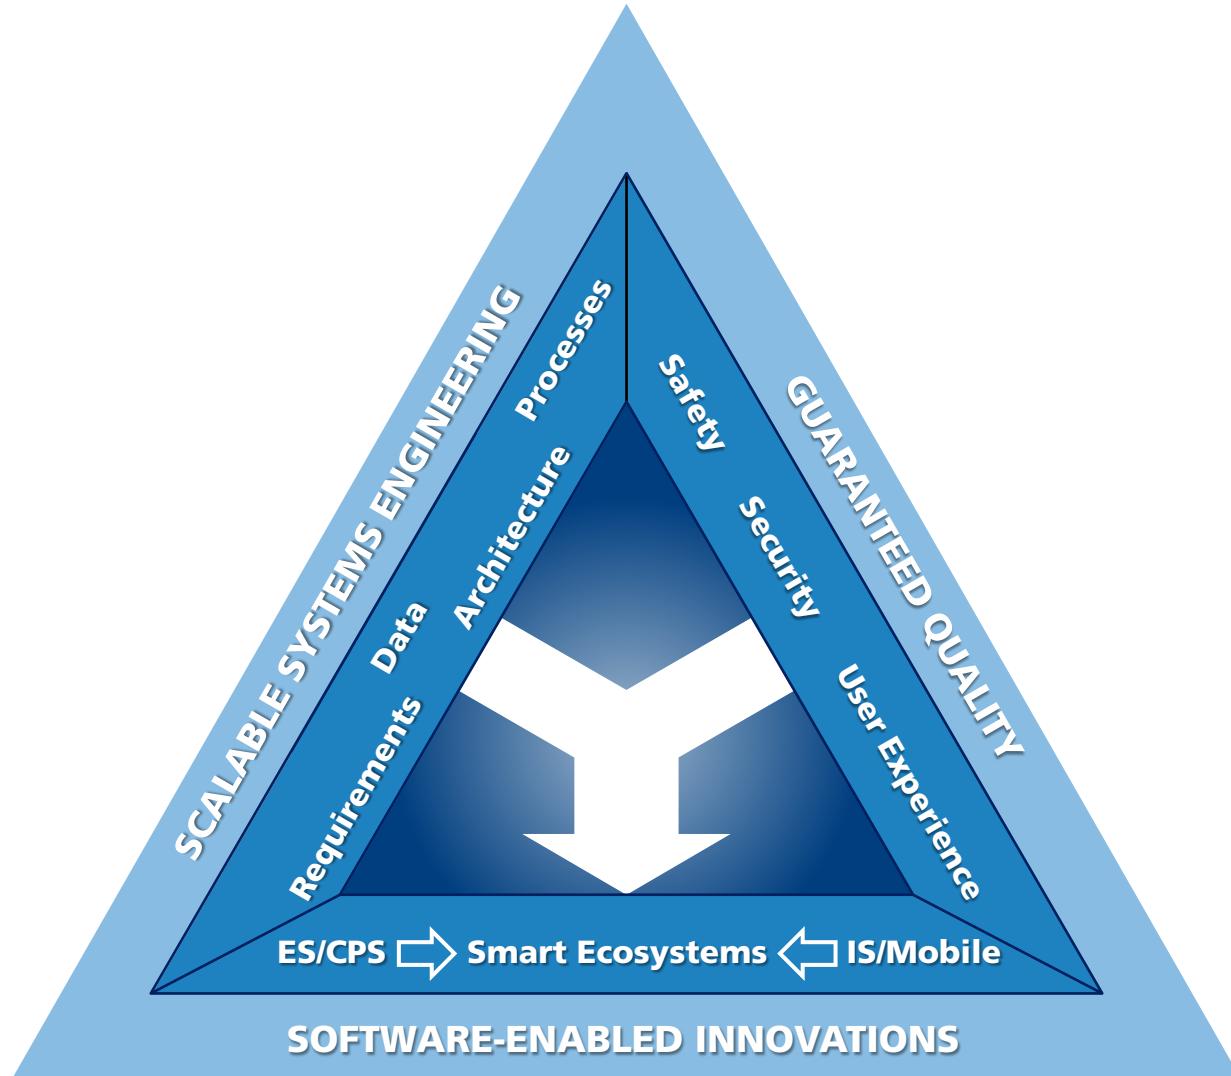


- Founded in 1996, headquartered in Kaiserslautern
- Over 155 full-time equivalents (FTEs)
- Our solutions can be scaled flexibly and are suitable for companies of any size

- Our most important business areas:
 - Autonomous & Cyber-Physical Systems
 - Digital Services
 - Defense & eGovernment



Our Competencies – for Your Benefit





Smart
Farming

Industry
4.0

Smart
Mobility

SMART
Ecosystems



Smart
Health

Smart
Energy



Smart
Rural
Areas

Smart X



Project Roots

Industry Cooperation



Cloud Models

Cloud Models

Based on
Domain Driven Design



Cross-Platform
Model-Driven Approach

Cloud Models

Replicated

Eventually Consistent

Built-in Conflict Management

Synchronization of High-Level
Service Operations



Multi-Level Transaction
Support



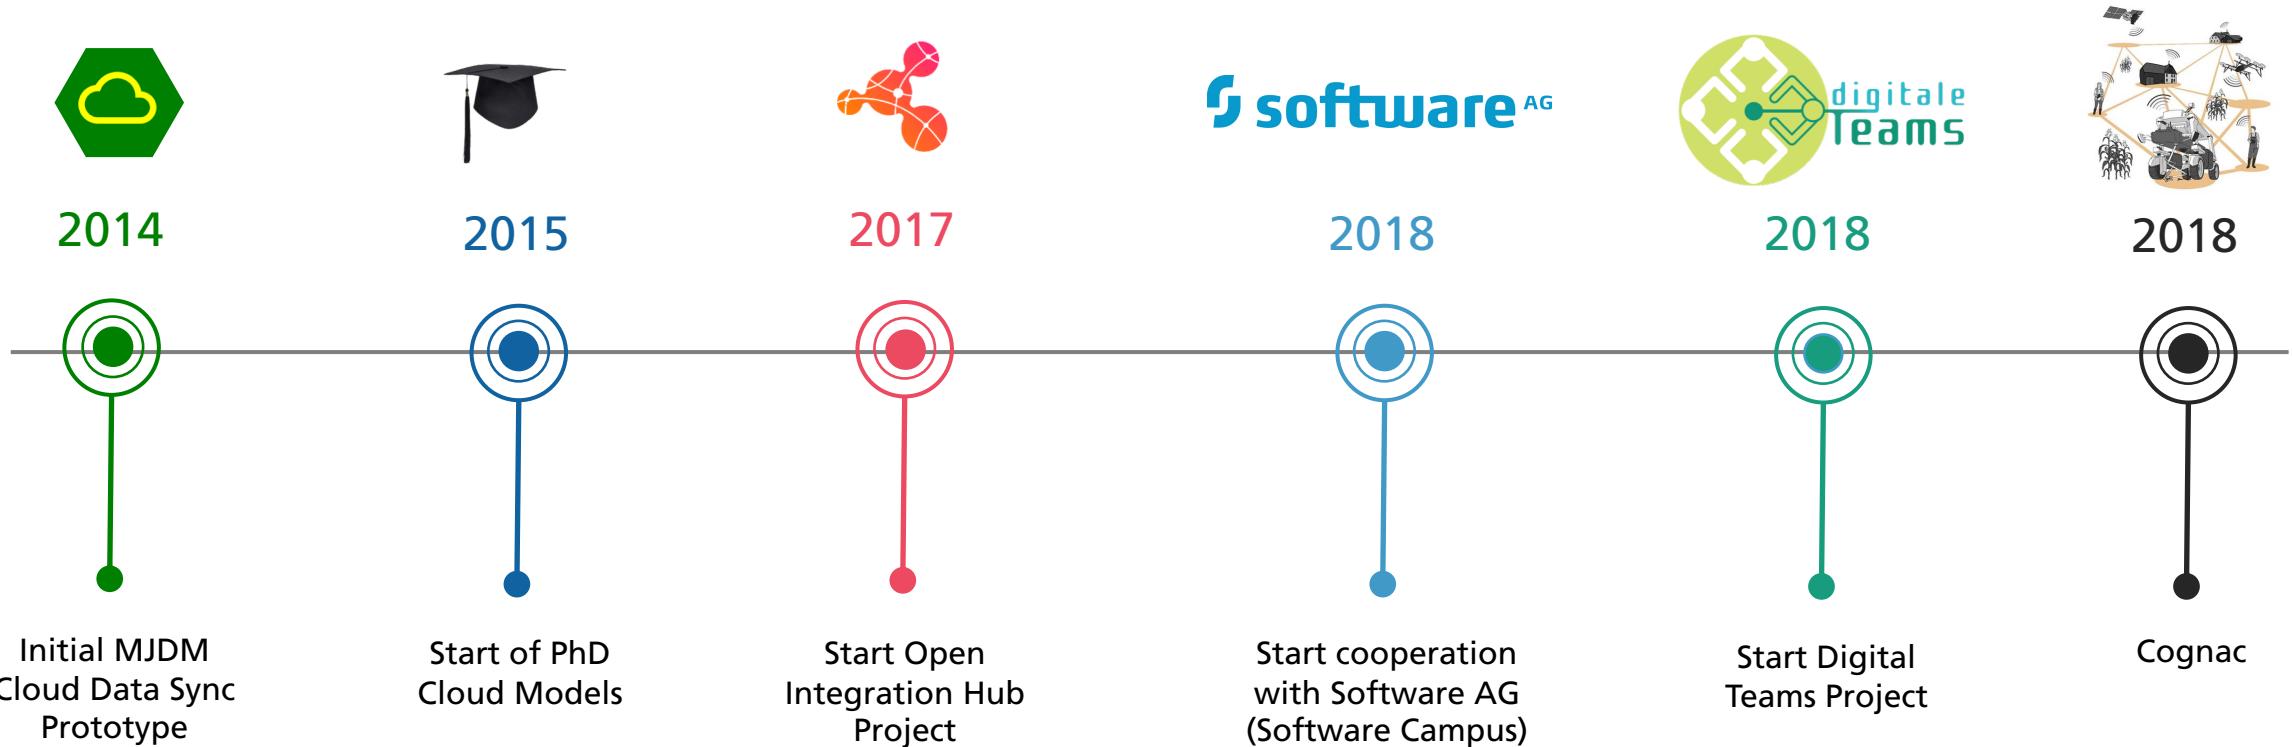


Federal Ministry
of Education
and Research



on the basis of a decision
by the German Bundestag

Cloud Models Funding



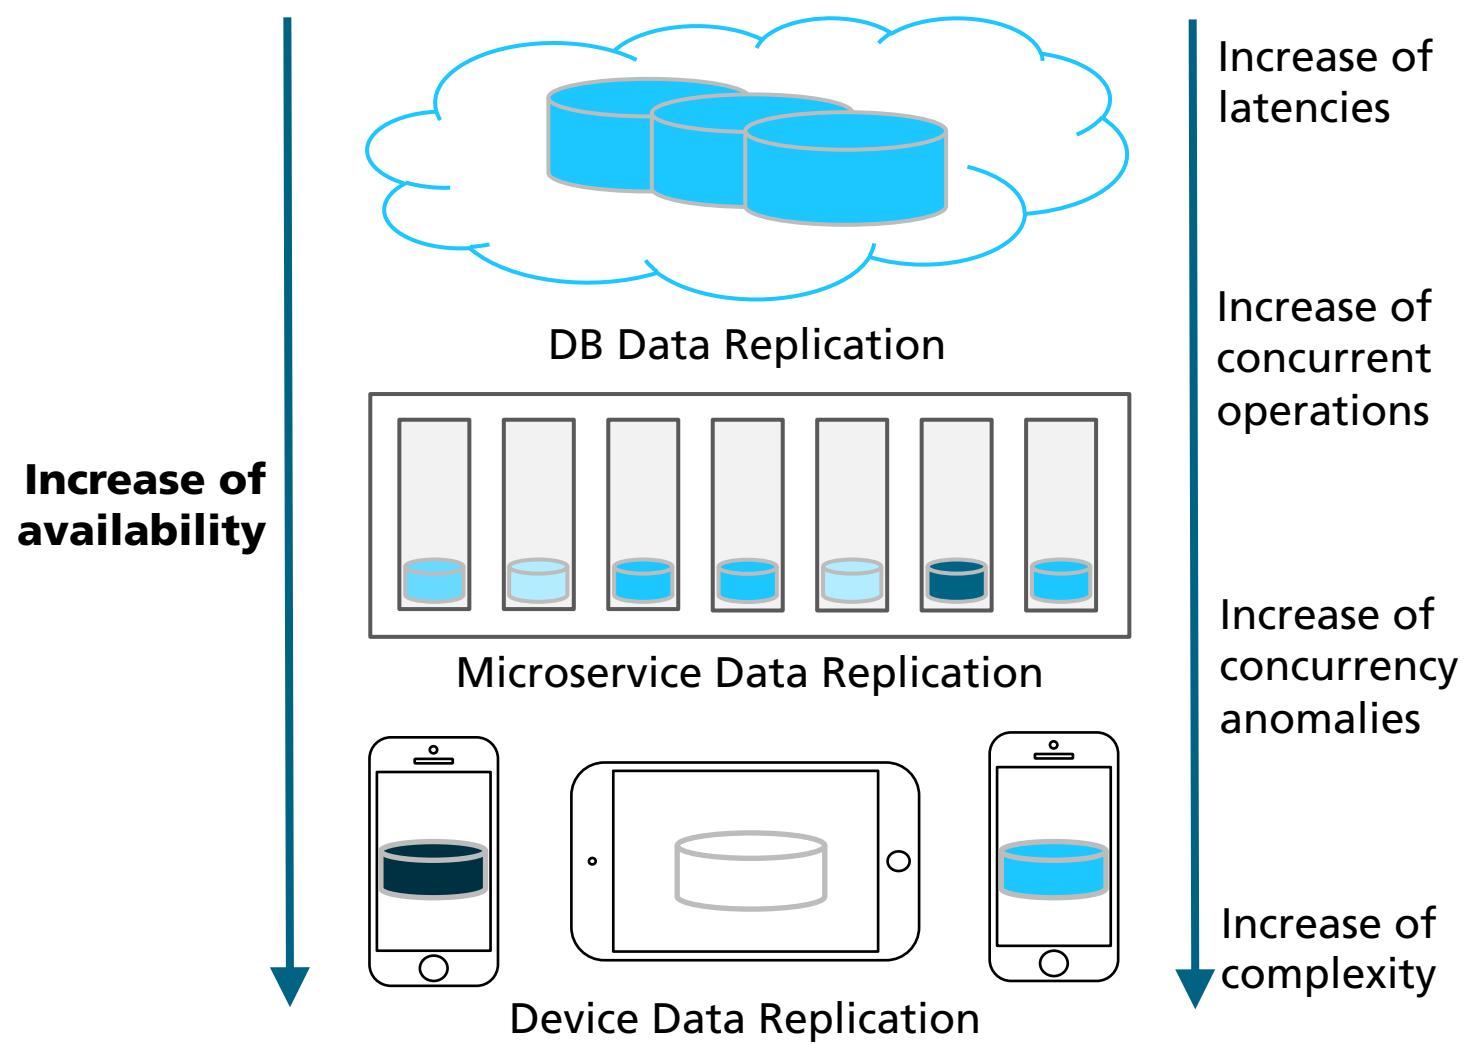
Key Requirements Replication in Ecosystems

- Offline capability of mobile apps
- Data synchronization of **core shared domain data** as platform service
- Support for standardized data models of **core shared domains**
 - No canonical data models

Key Requirements Replication in Ecosystems

- Modularity
- Horizontal Scalability
- Loose coupling of apps and services of the ecosystem
- High availability of apps and platform services, in particular data synchronization services
- Resilience of apps and platform services, in particular data synchronization services

Relevance for Modern Software Architectures



Goals for Cloud Models in Open Integration Hub

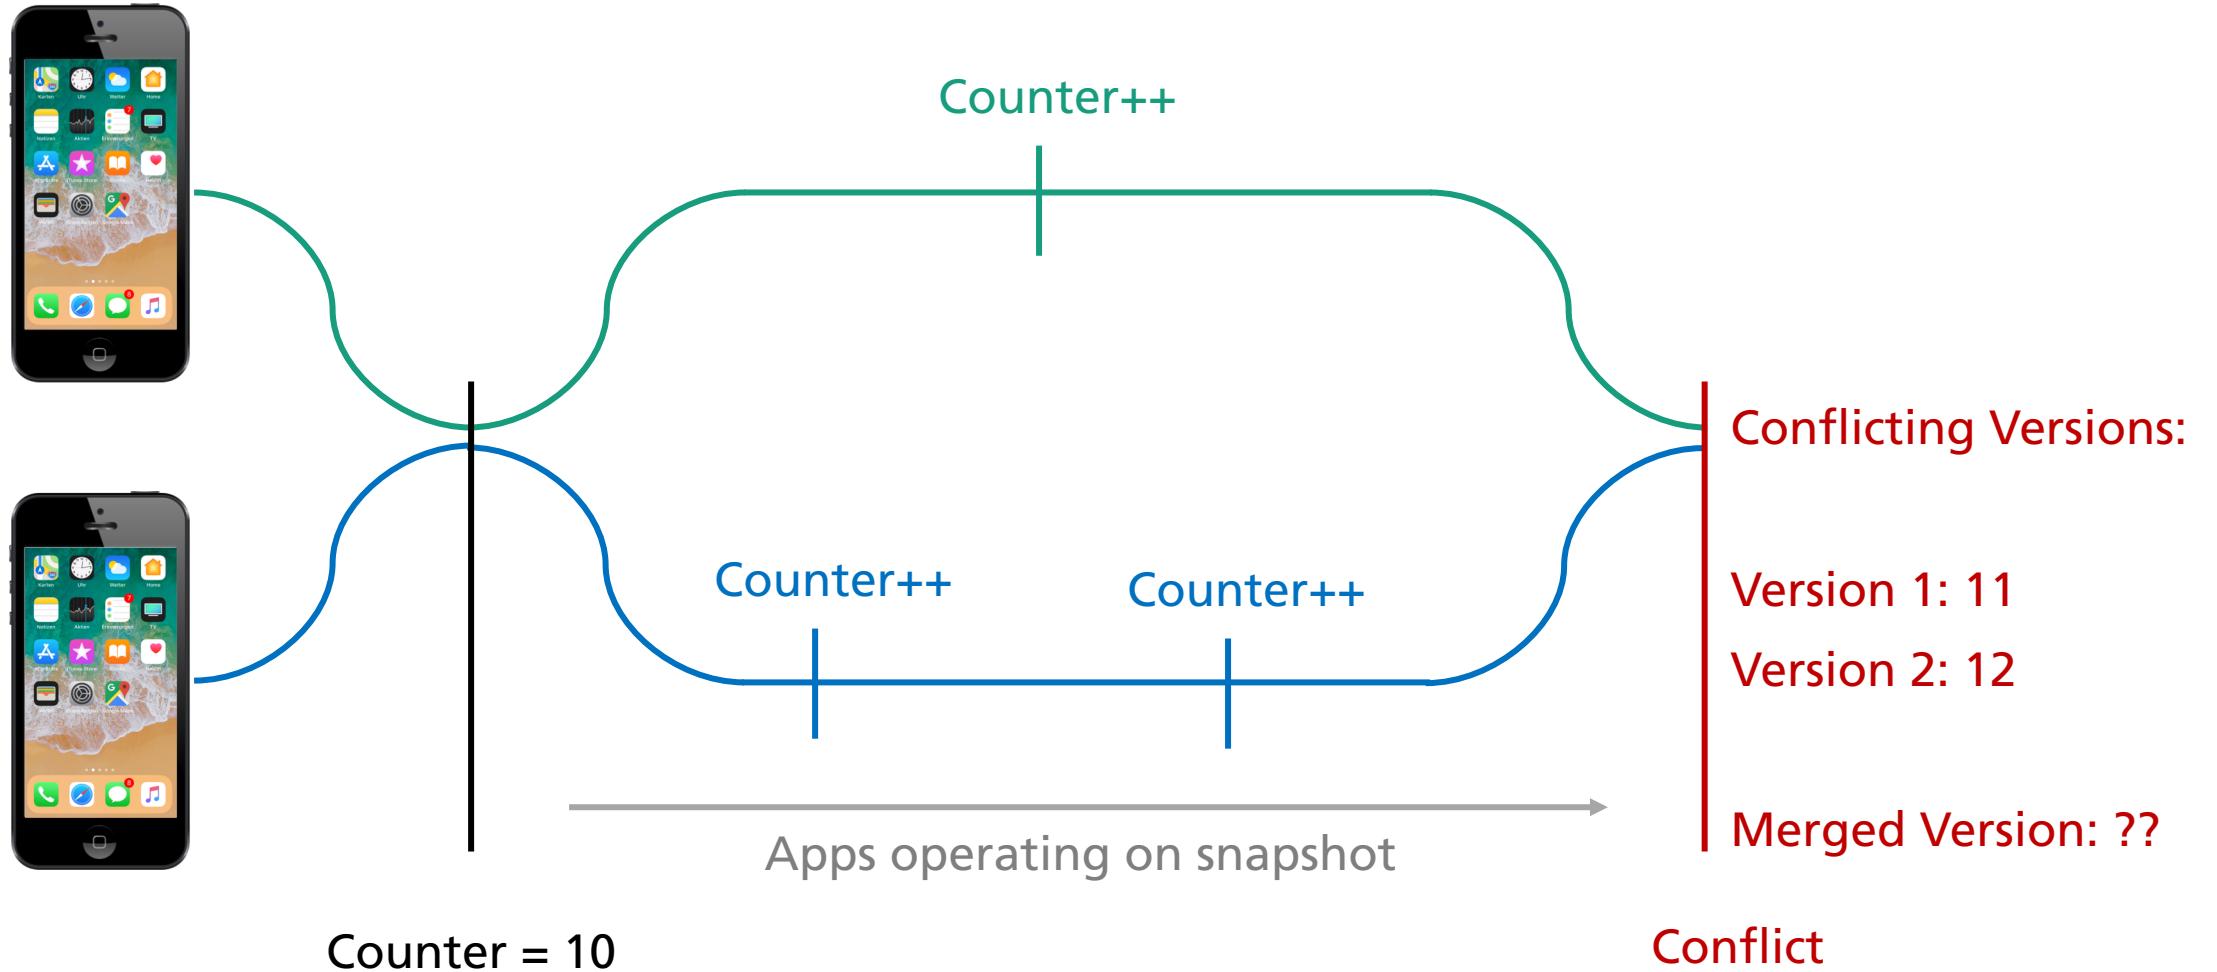
- Compare IESE approach with OIH approach
- Make small evaluation
- Publish results at scientific conferences
 - Wissenschaftliche Verwertung!

Problems

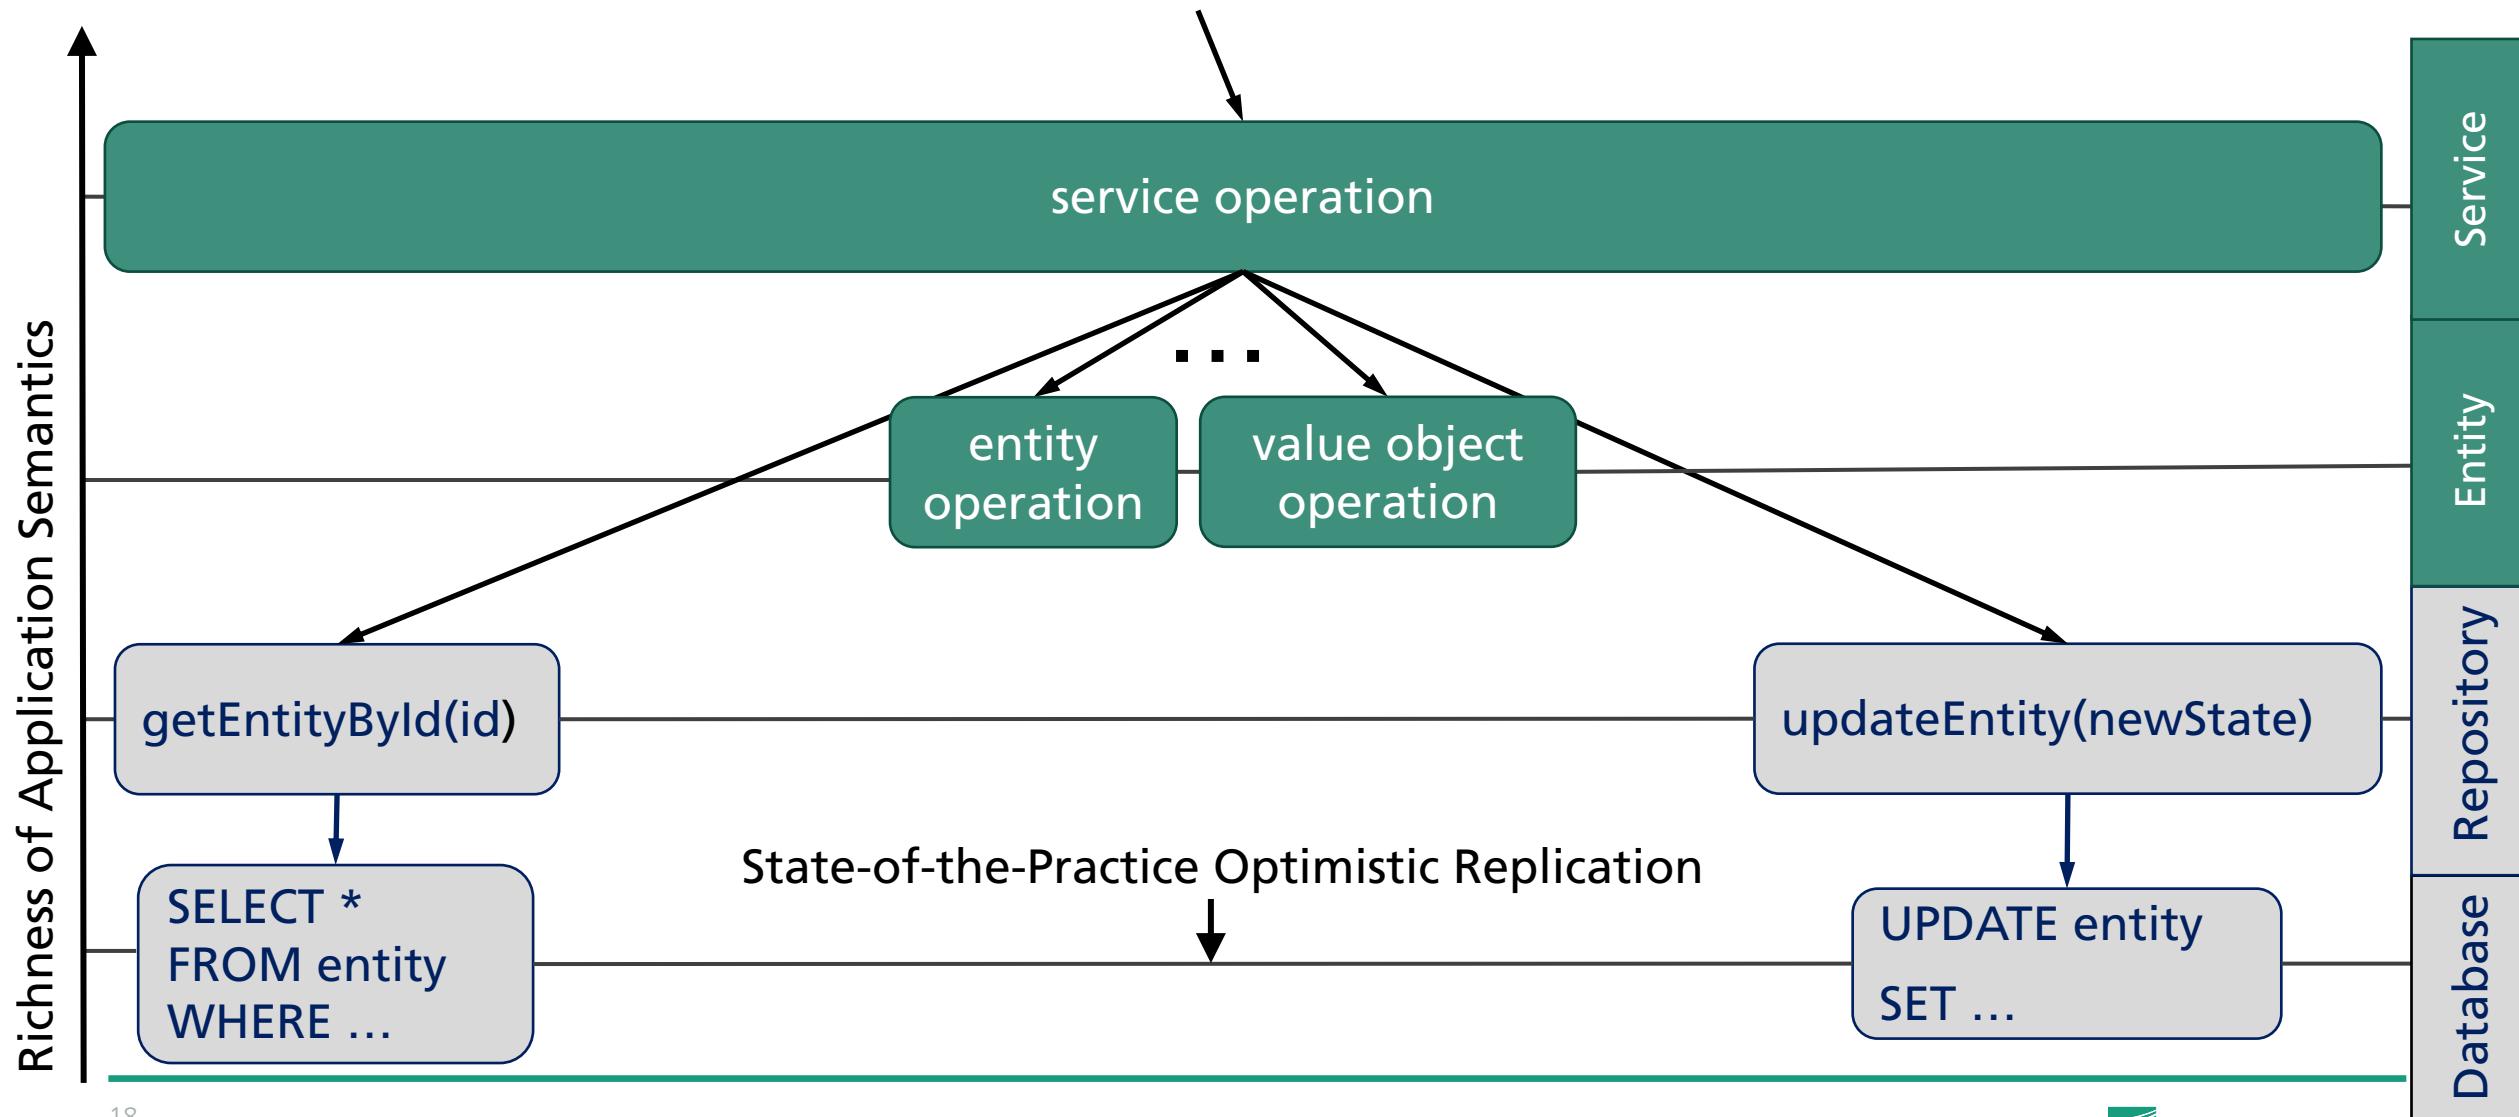
Eventual Consistency

- **Updates are propagated asynchronously (high latencies possible)**
 - Drivers: Availability, Scalability, Loose Coupling, Resilience
- Stale Data
- Conflicts
- **Merge of conflicting versions**
 - Unintuitive and error prone
 - Huge source of human error

Shared Counter Example



State-of-the-Practice Optimistic Replication



Issues in Practice

- Conflicts that cannot be resolved correctly
 - Ultimately resulting in lost updates
- Conflicts on the semantic level that cannot be detected syntactically
 - Ultimately resulting in constraint violations
- Pseudo-conflicts detected syntactically that do not conflict on the semantic level but have to be reconciled
 - Thereby adversely impacting performance

Recap Concurrency Control in Relational DBs

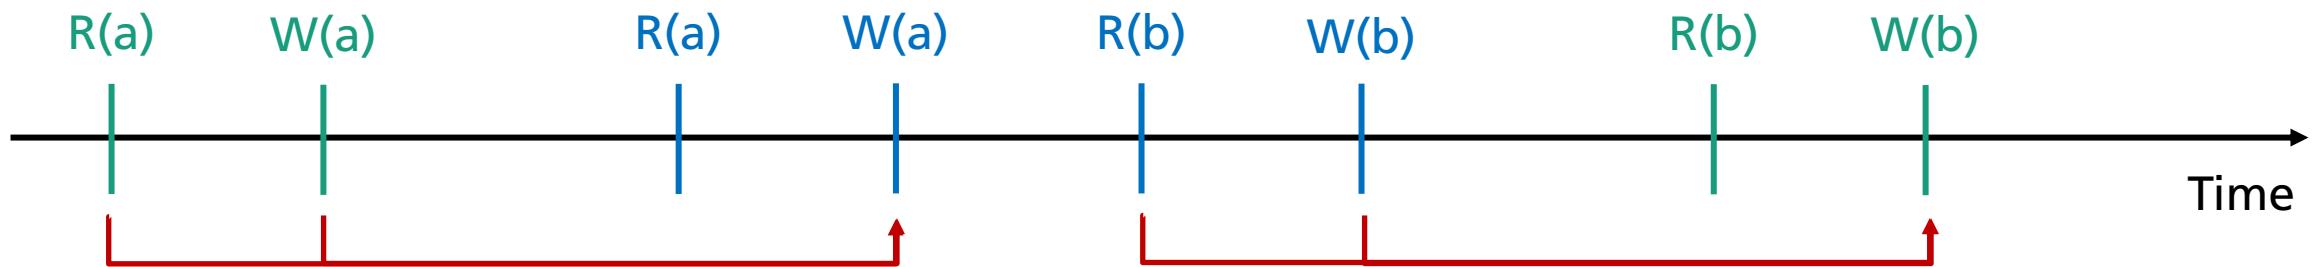
Transaction Interleaving Approach - Conflict Serializability

- Serial execution of transactions prevents occurrence of concurrency anomalies
- Conflict Serializability
 - $R(x)$ is compatible with $R(x)$
 - $R(x)$ is in **conflict** with $W(x)$
 - $W(x)$ is in **conflict** with $W(x)$
- A schedule of concurrent transactions is conflict-serializable iff the conflict serialization graph is acyclic

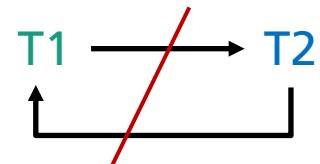
Example Conflict Serializability

- A schedule of concurrent transactions is conflict-serializable iff the conflict graph is acyclic and compatible with the execution order of the conflicting operations

Transactions T_1 , T_2 :



Conflict Graph:

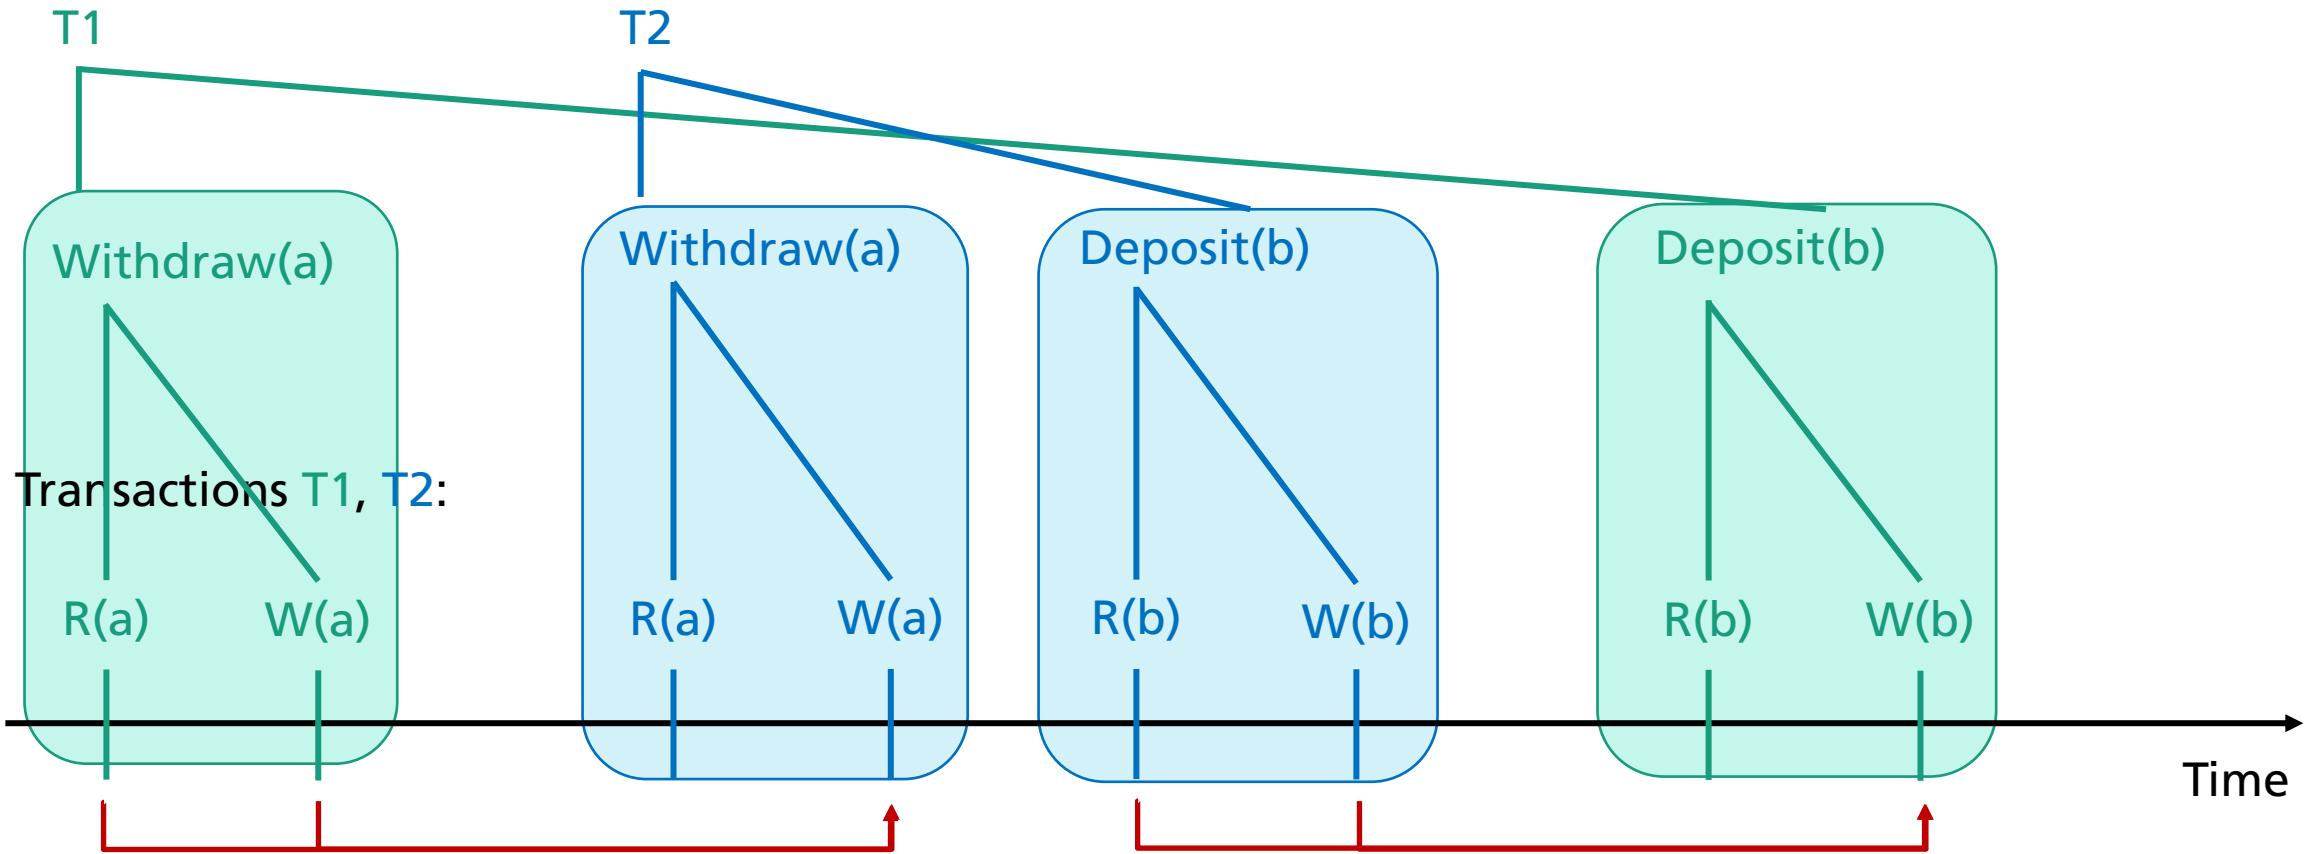


Conflict Serialization Graph is cyclic

→ No conflict serializability

→ Schedule would be rejected

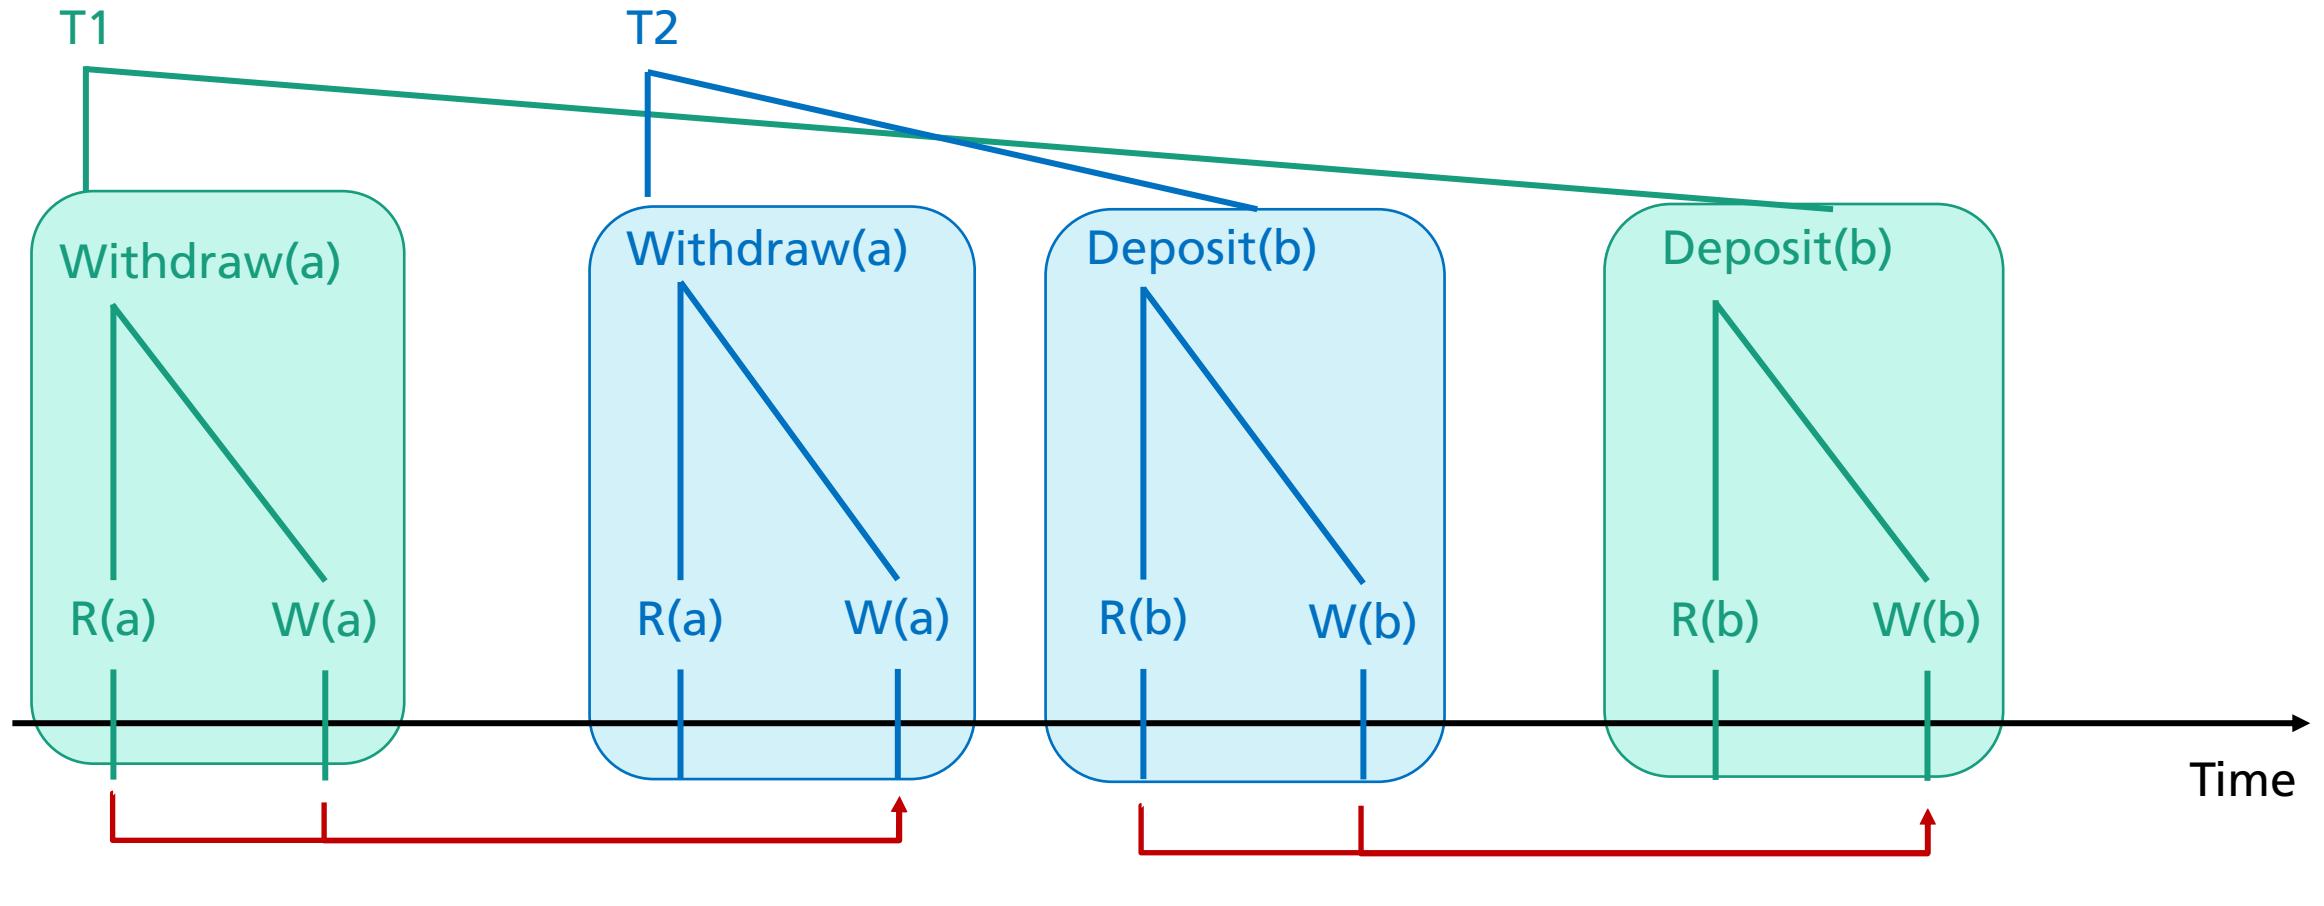
Banking Example



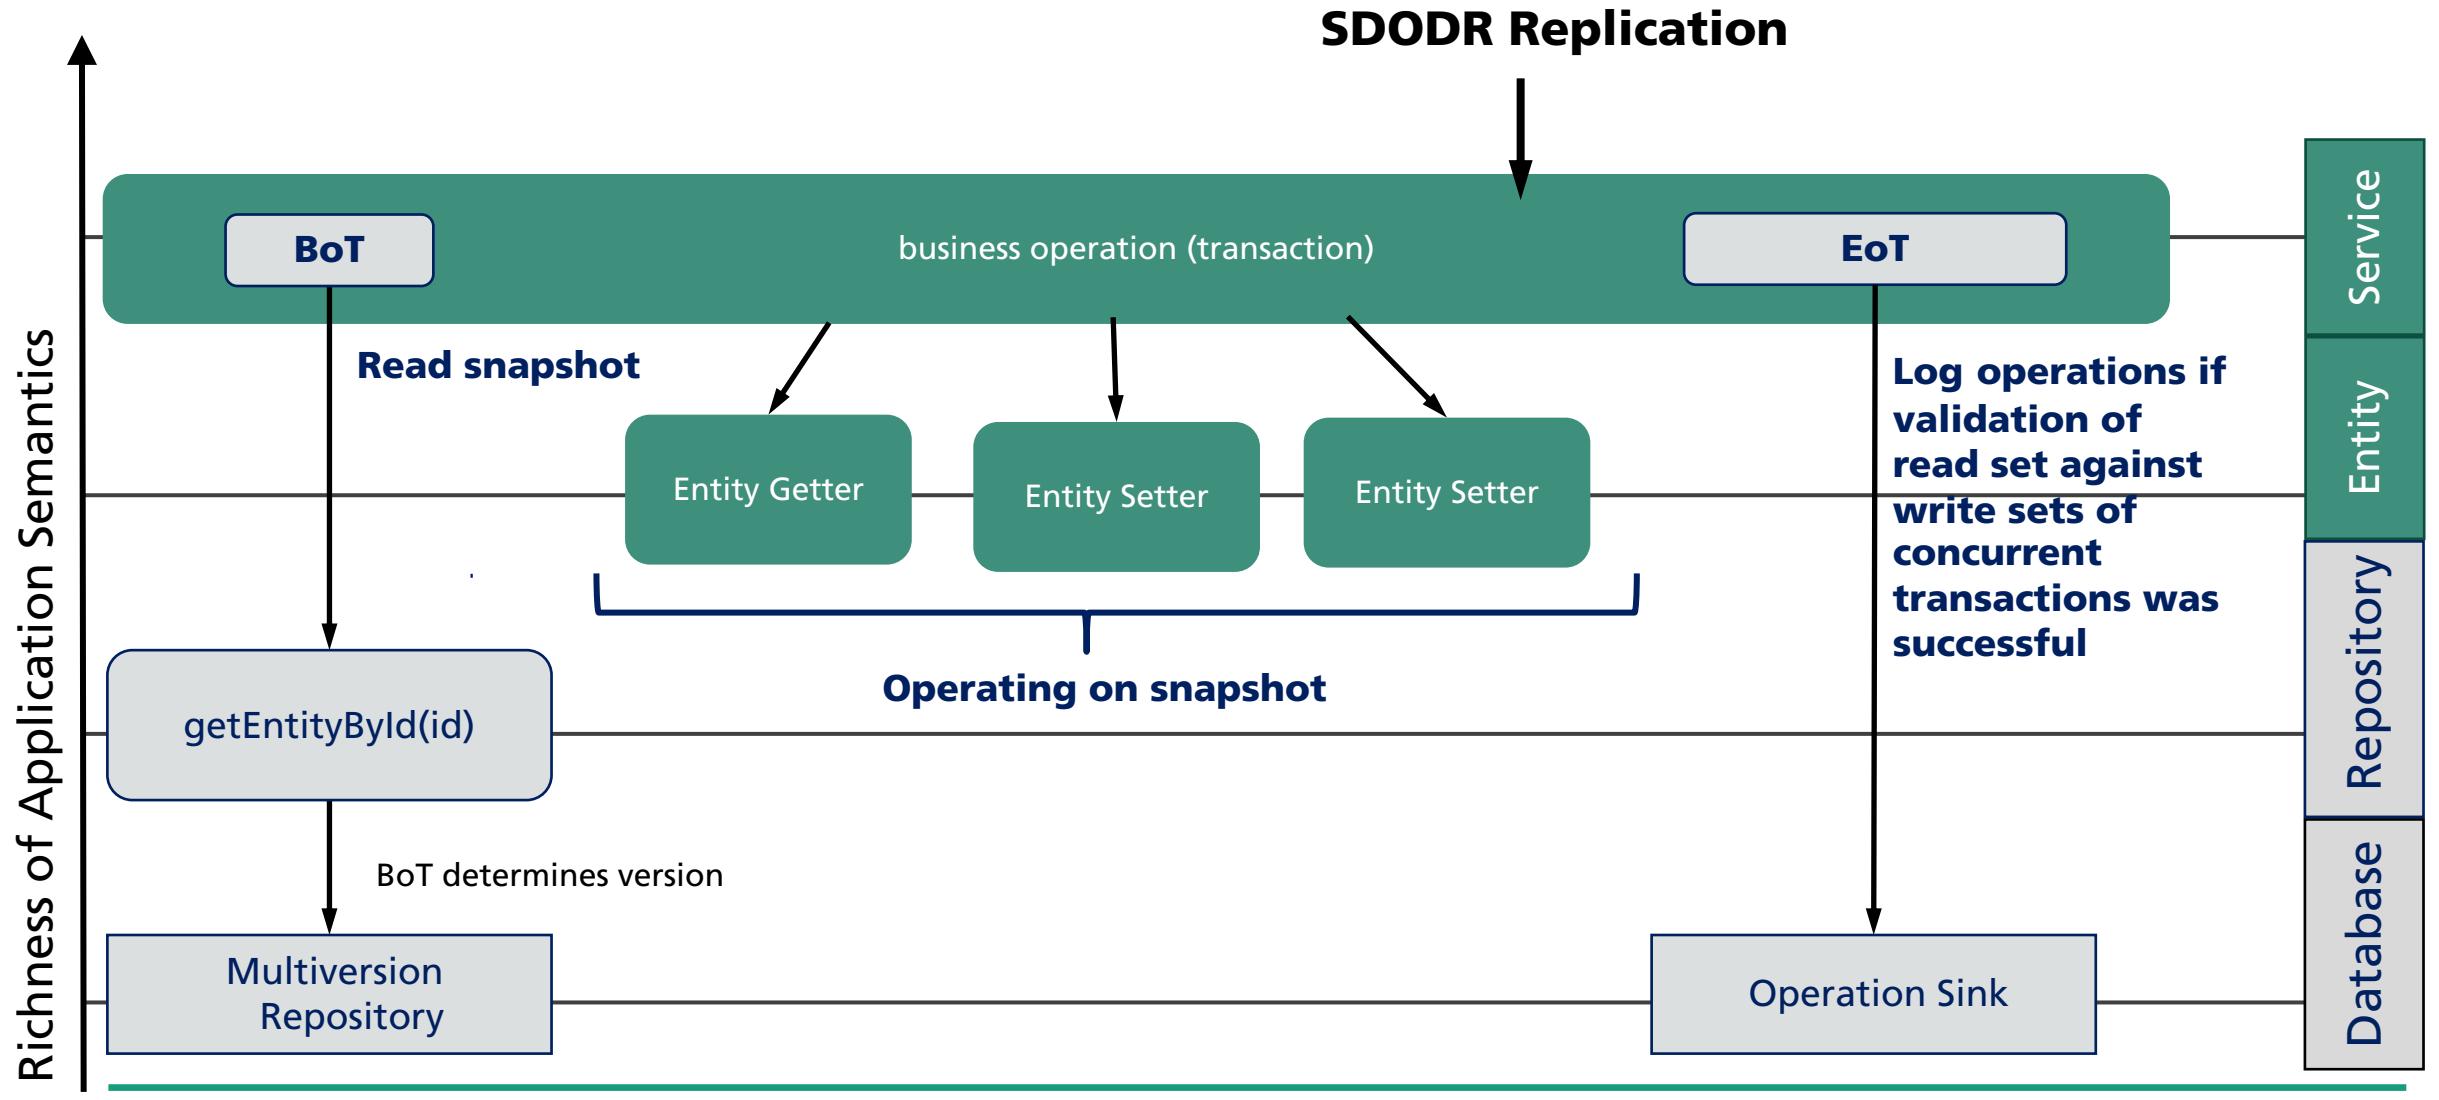
Multilevel Transactions

- **Key Idea: Exploit the semantics of operations in level-specific conflict relations that reflect the commutativity or compatibility of operations to increase concurrency**
- Transactions are decomposed into operations and the operations again into sub-operations on multiple levels
 - Transactions, Business operations, Low-level read and write operations
- At each level a conflict relationship is defined
 - read-write conflicts and write-write conflicts on the same data item conflict at the lowest level
 - Non-commutative operations are conflicting on the level of business operations
- If at each level the conflict serialization graph is acyclic then the multilevel schedule is in total multilevel serializable

Banking Example



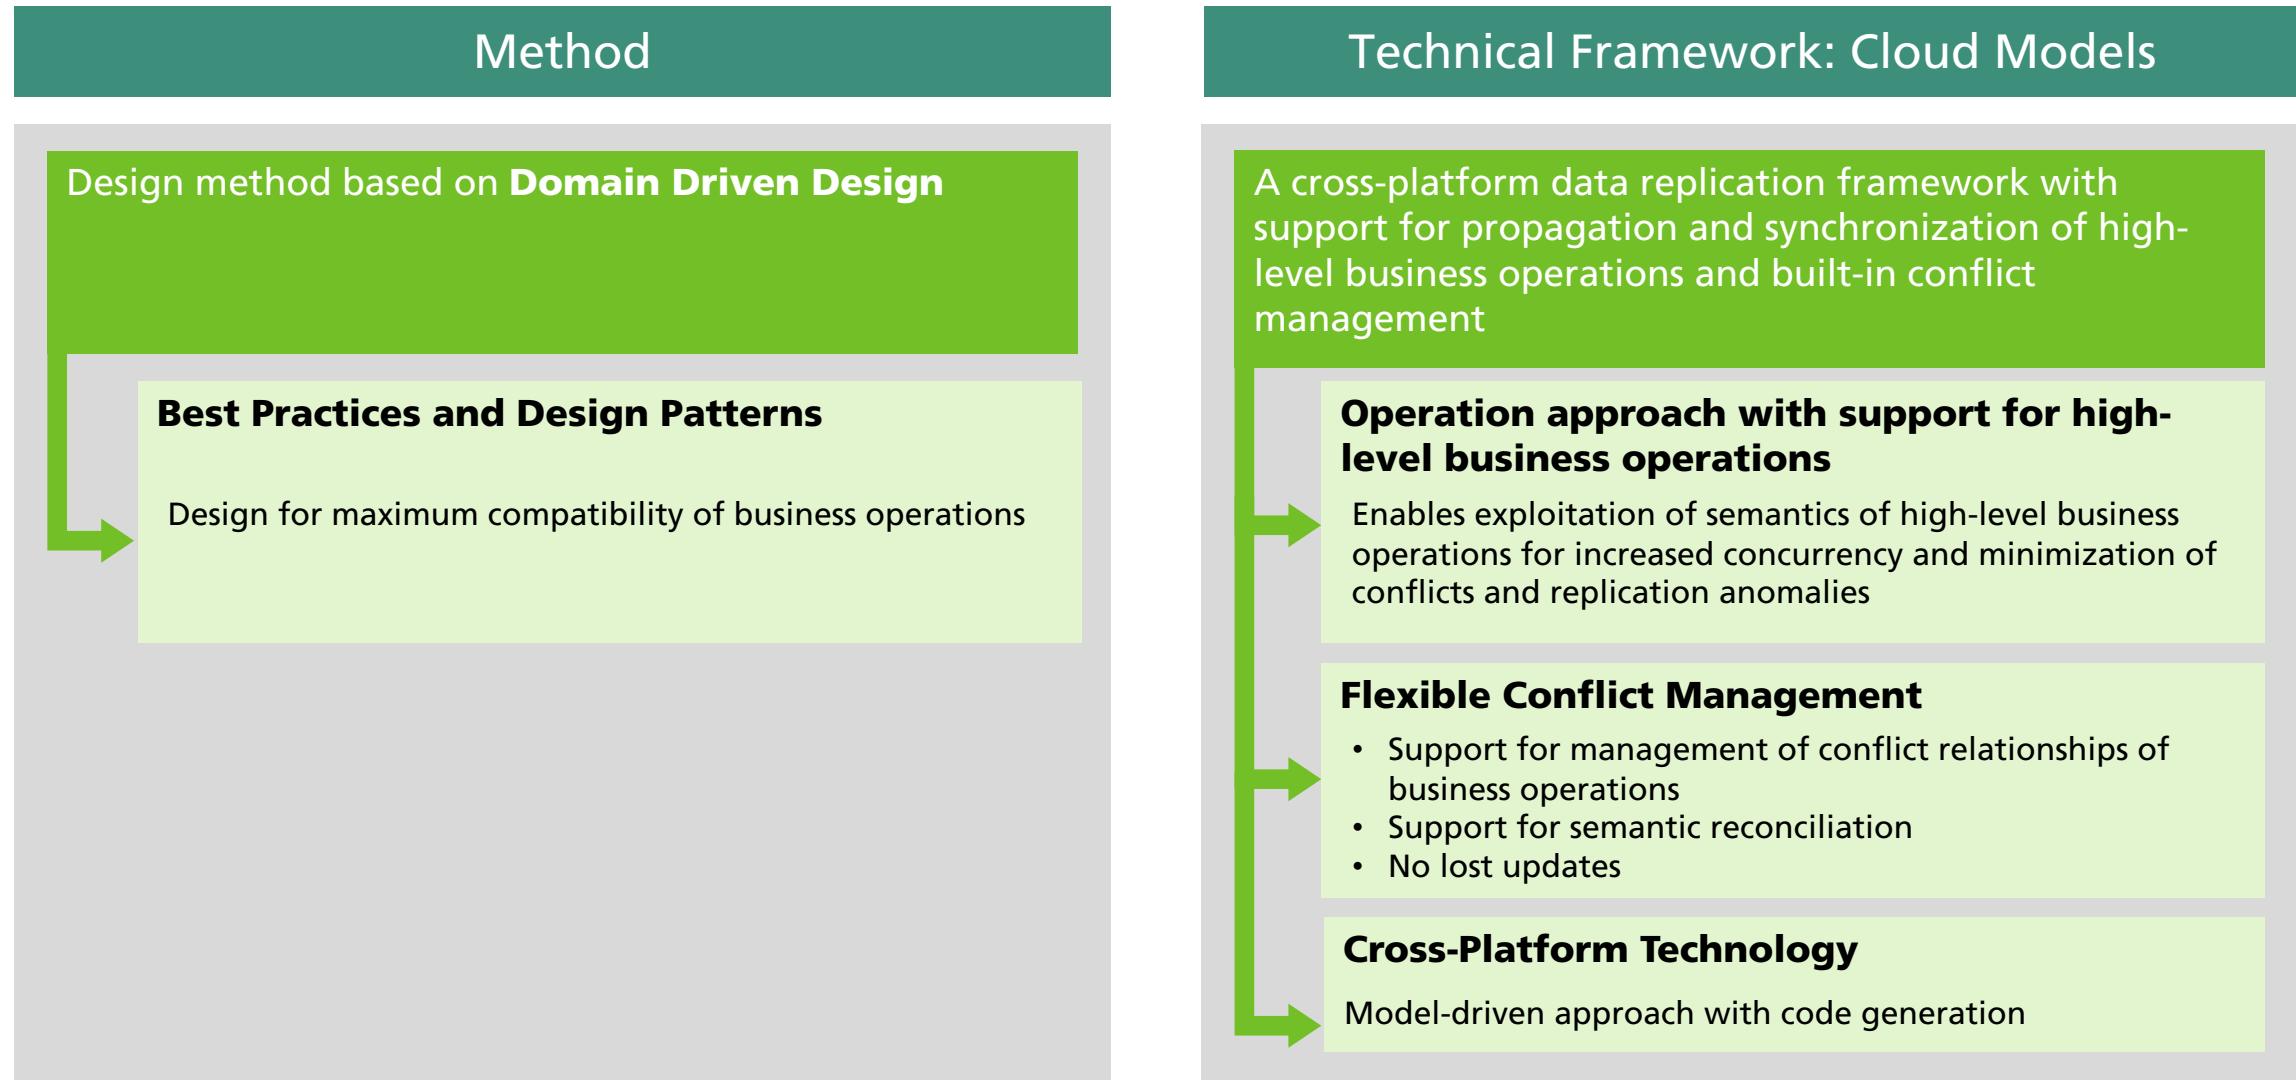
Level of SDODR Replication



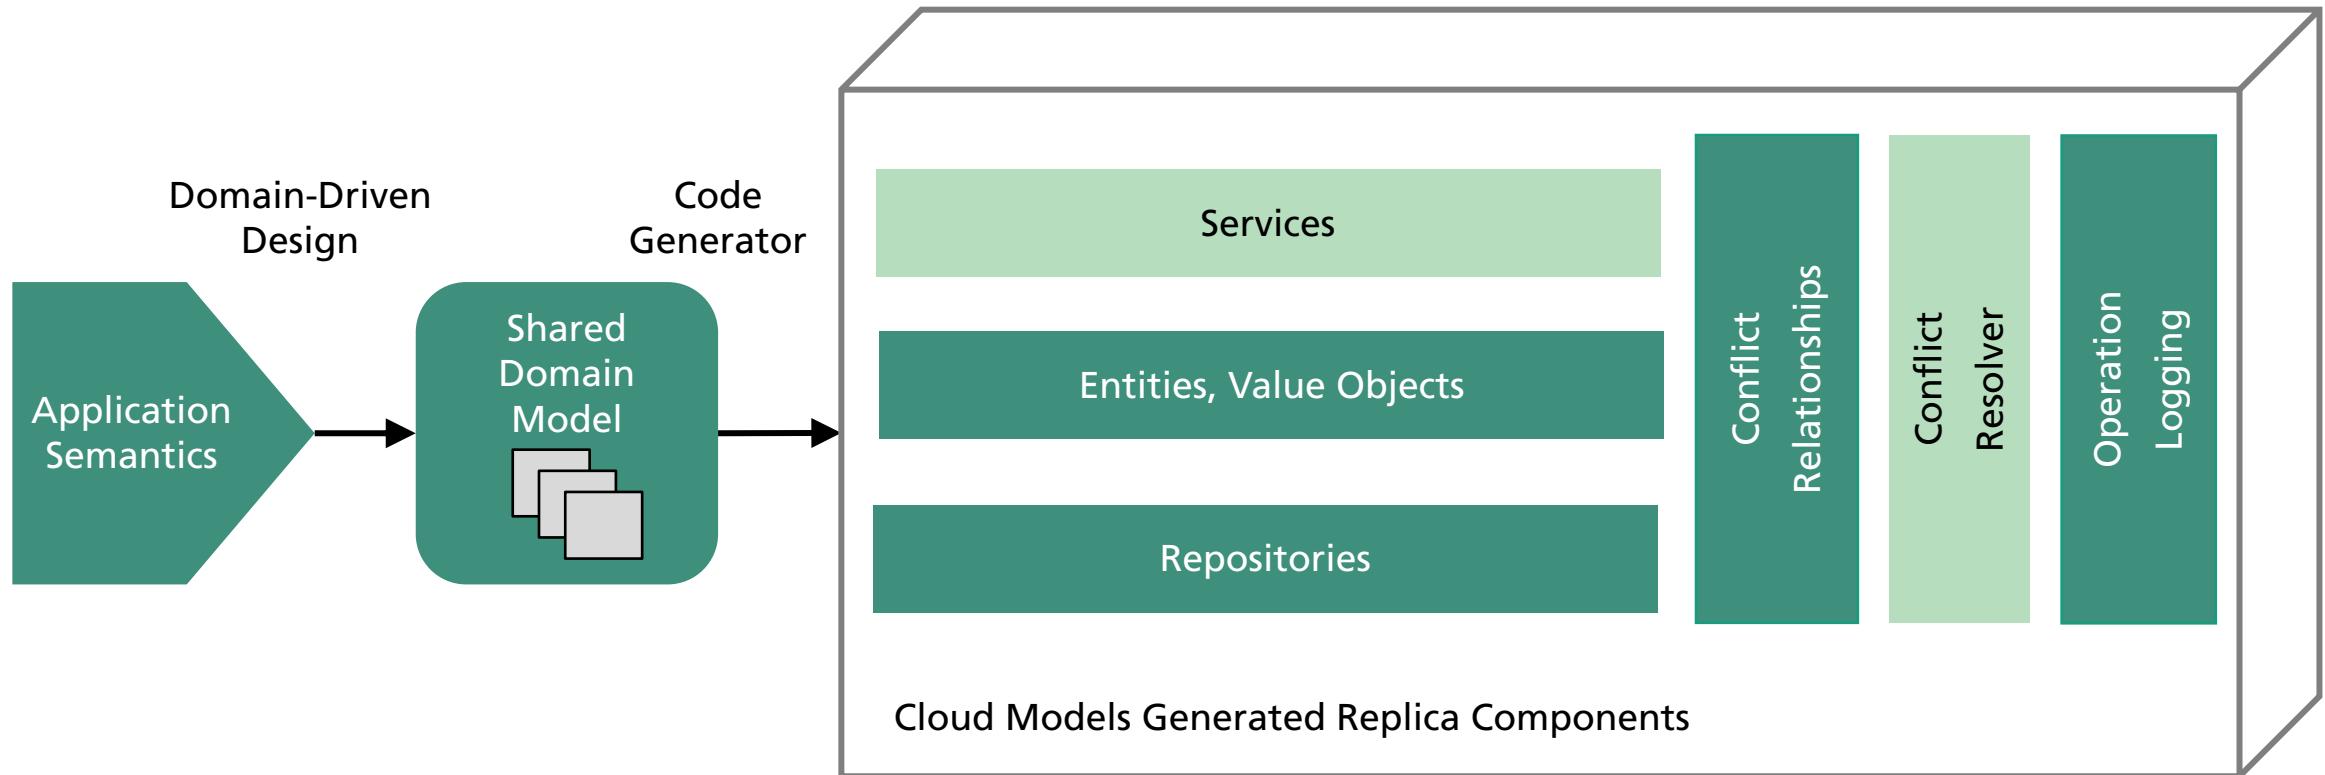
Idea

Exploit application semantics to increase
number of operations that can run
concurrently and conflict-free on different
replication nodes

Semantics-Driven Optimistic Data Replication (SDODR)



How to provide a framework that can exploit application semantics for concurrency control?

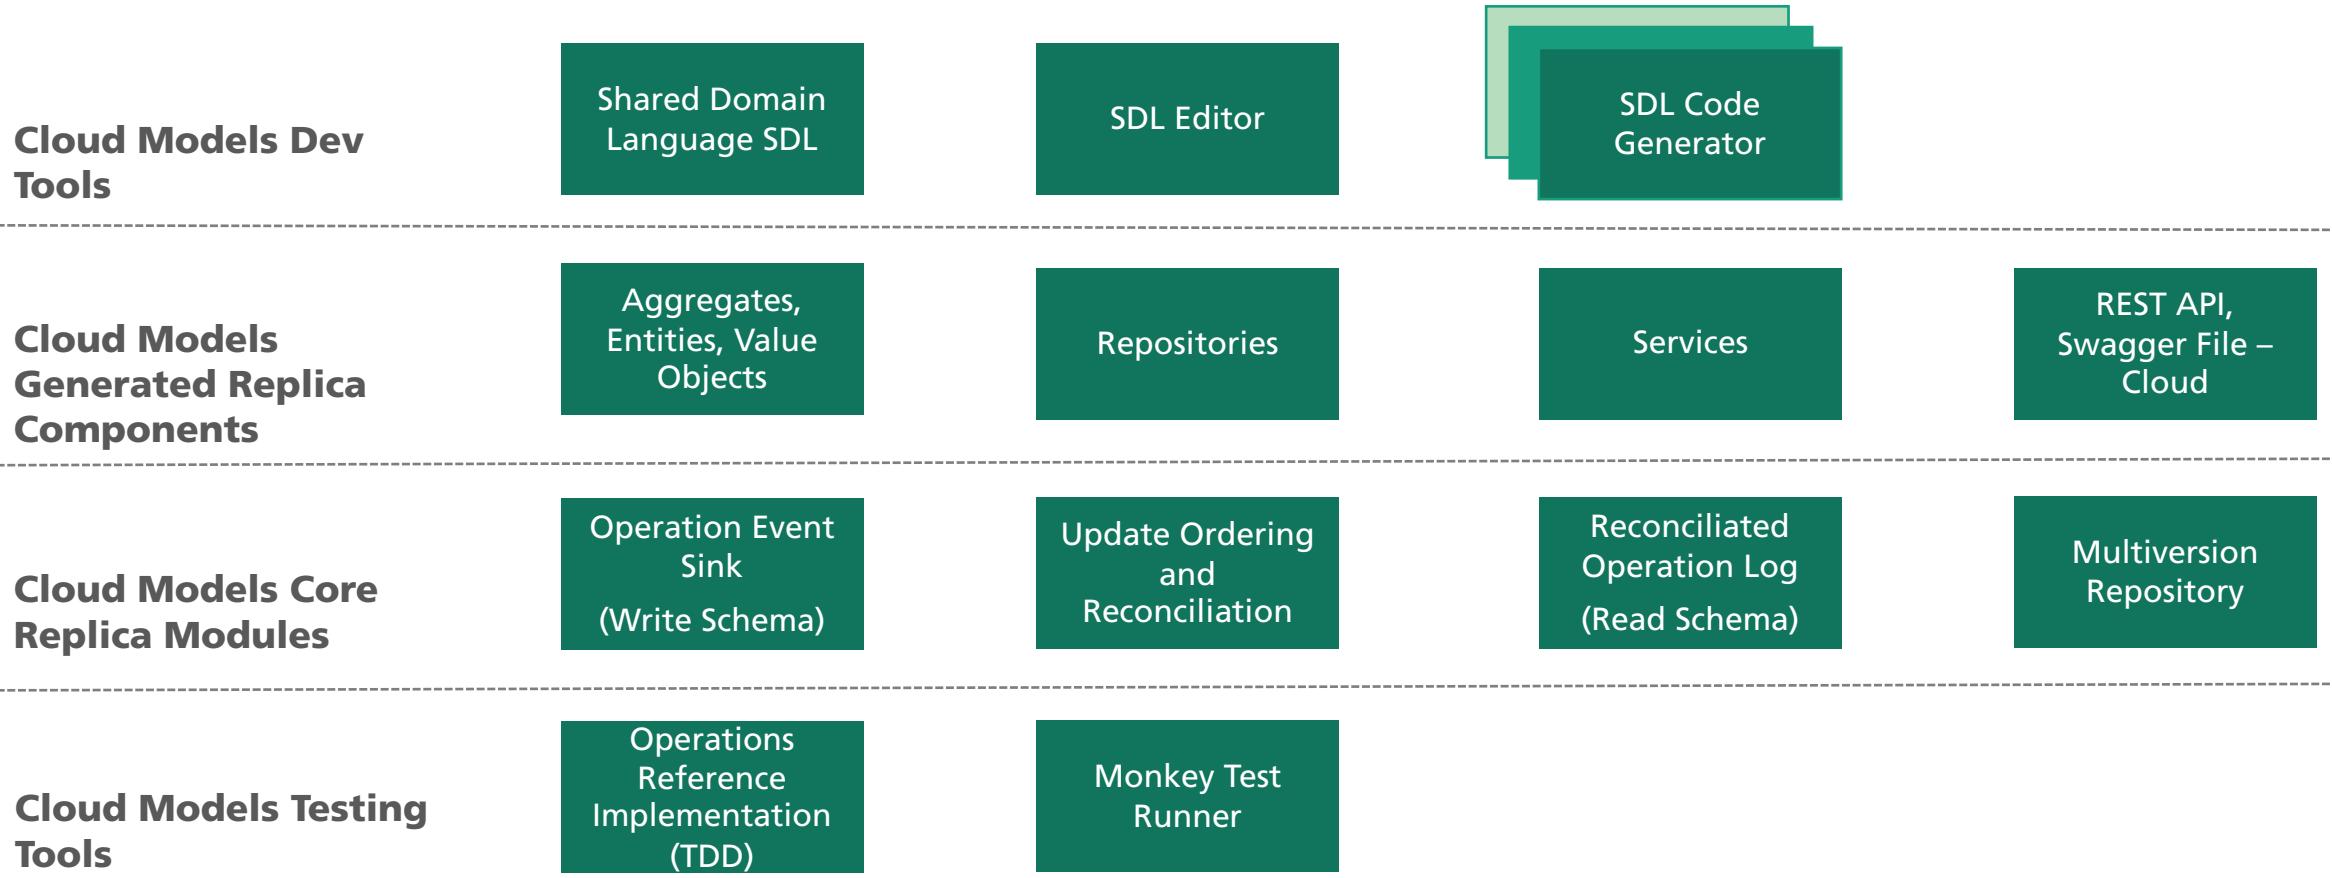


Shared Domain Model Example

```
Namespace de.fraunhofer.iese.oih  
Shared Domain Contacts  
Version 1.0  
  
Entity Contact {  
    person: Person  
    organization: Organization  
    phoneNumbers: List of PhoneNumber  
}  
Entity Person {  
    firstName: String  
    lastName: String  
}  
Entity Organization {  
    name: String  
    address: Address  
}  
Value Object Address {  
    street: Street  
    city: City  
}  
ValueObject PhoneNumber {  
...}
```

```
...  
Service ContactsService {  
  
    Operation deletePhoneNumber(Contact contact, Integer phoneNumberIndex): Contact  
        conflicts with changePositionOfPhoneNumber  
        resolve with resolveConflictOnListOfPhoneNumbers  
  
    Operation changePositionOfPhoneNumber(Contact contact, Integer newPhoneNumberIndex): Contact  
        conflicts with deletePhoneNumber  
        resolve with resolveConflictOnListOfPhoneNumbers  
  
    ConflictHandler resolveConflictOnListOfPhoneNumbers  
}  
...
```

Software Architecture – Major Building Blocks

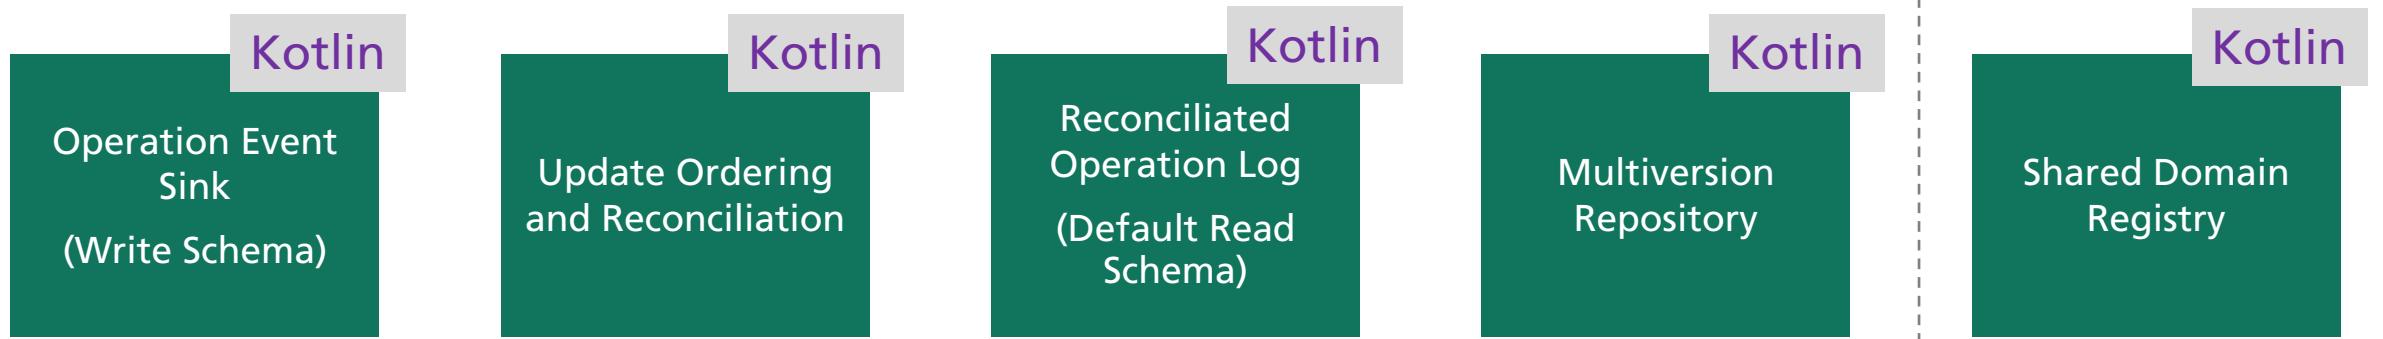


Cloud Models Dev Tools



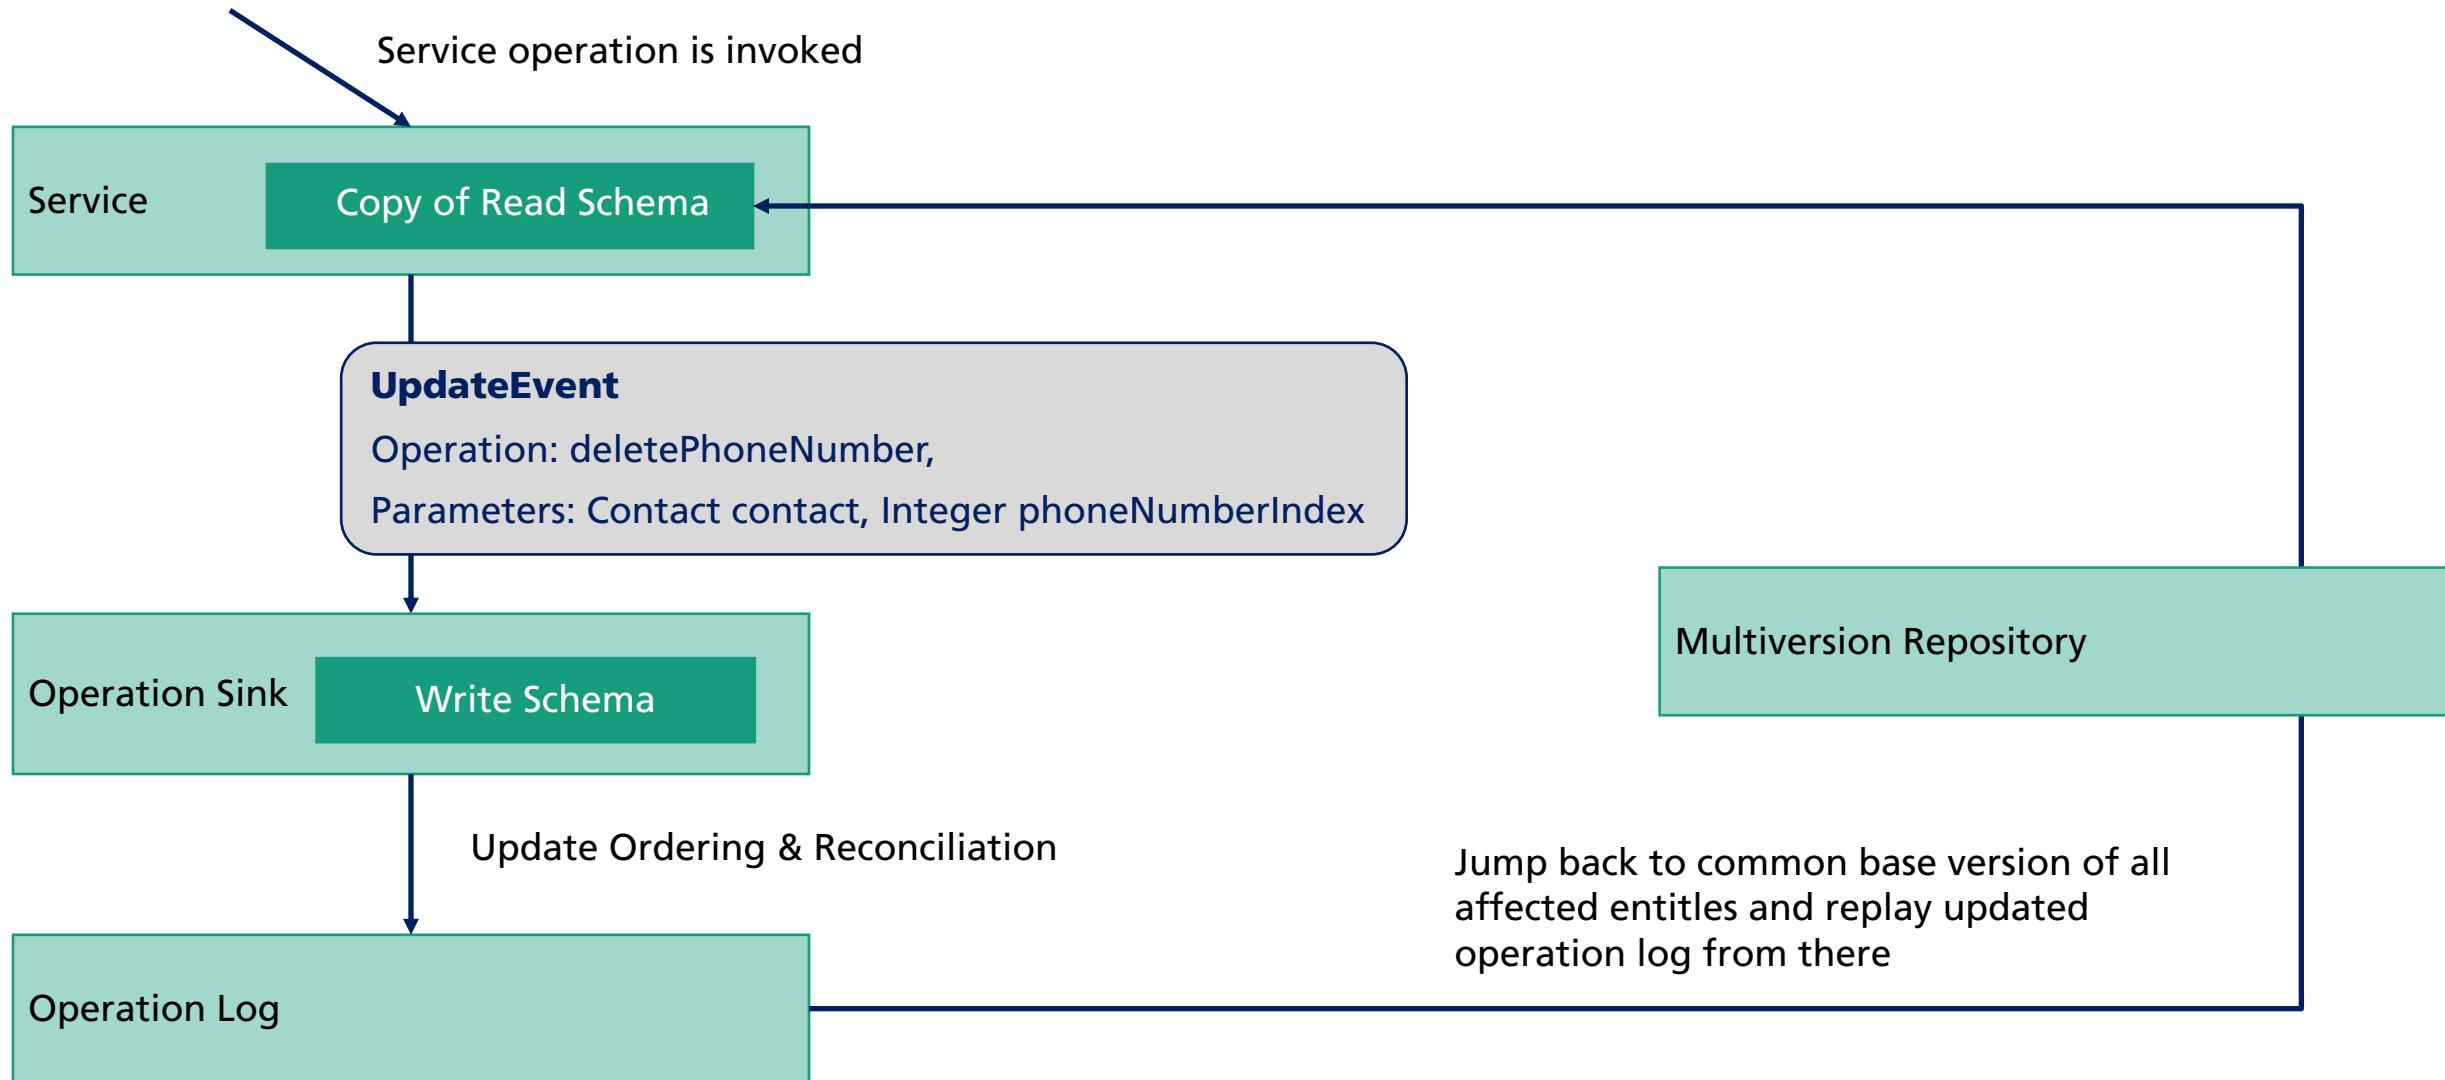
MVP

Cloud Models Core Modules (Cloud and Local)

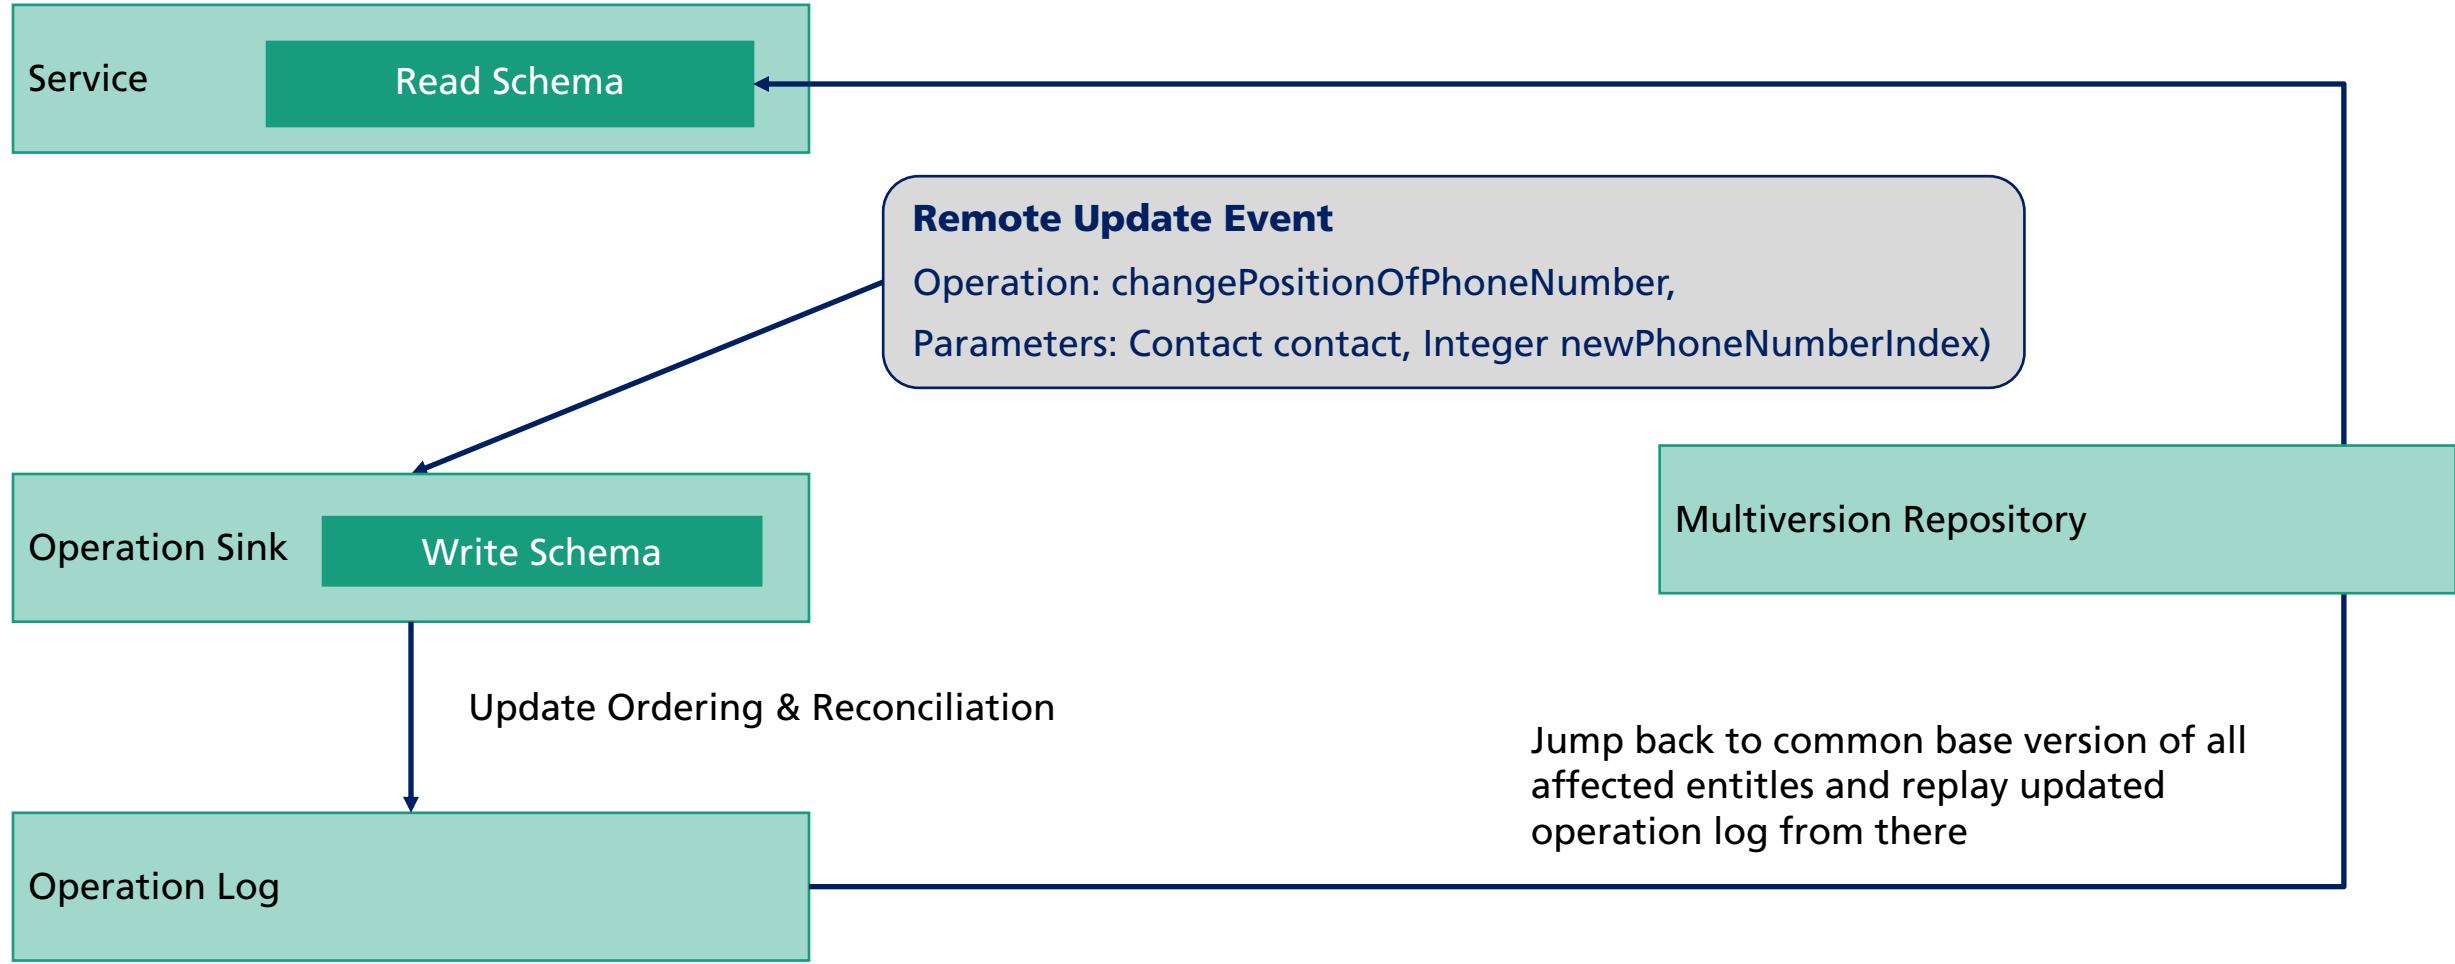


MVP

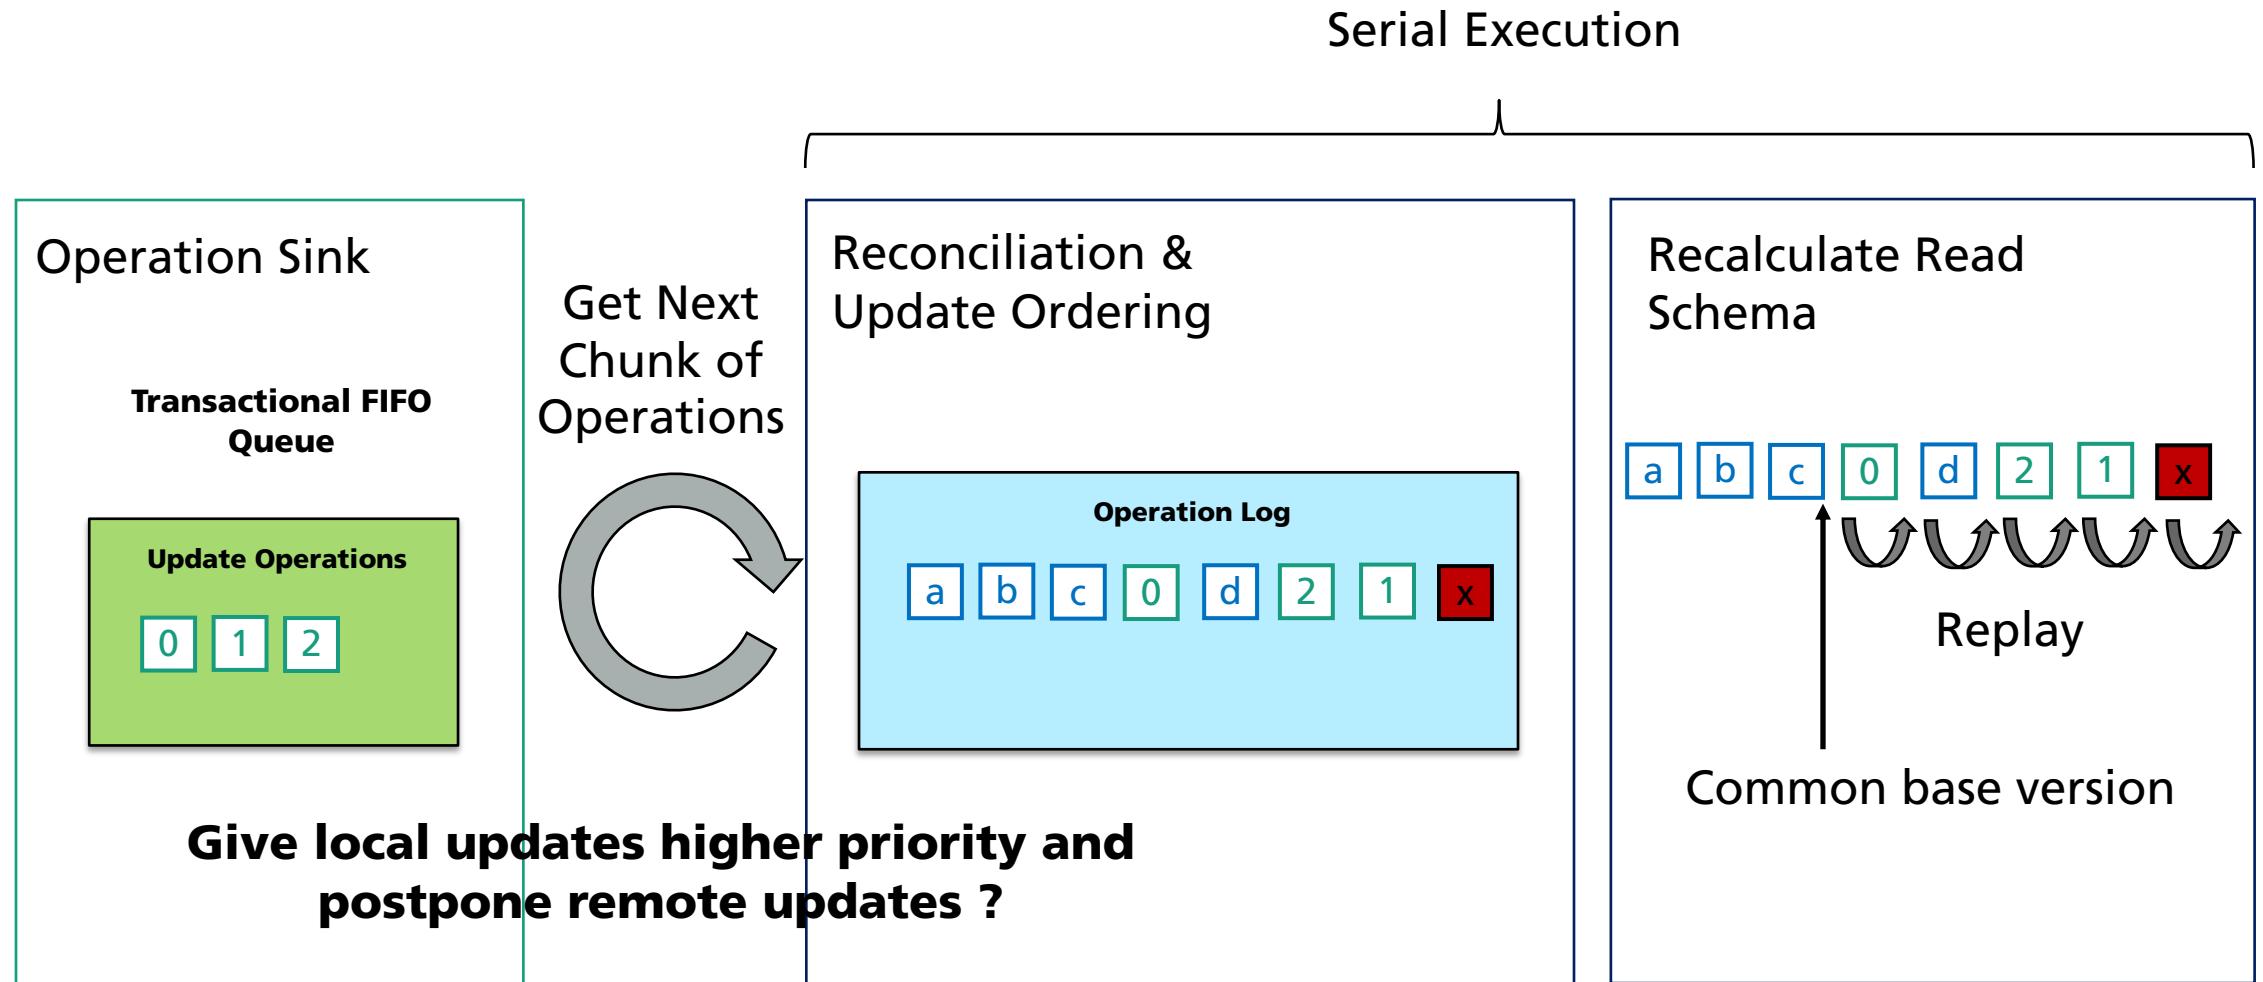
Distributed CQRS With Local Update Event



Distributed CQRS with Remote Updates



CQRS Performance Considerations (App / Microservice)



Operation Event Sink and Global Operation Log in the Cloud

- Transactional Queue System (e.g. Kafka) shared by different services (that scales well horizontally)
 - One “operation event sink” queue **per** “Master Data Model” (**shared by all cloud replicas**)
 - Updates can be “submitted” via generated REST API (similar to the one I proposed)
 - Example: POST to /contacts/update-events
 - Updates can be polled via the same generated REST API (similar to the one I proposed)
 - Example: GET /contacts/update-events
 - services maintain their own pointer in the transactional log (compare to how LinkedIn does it, Kappa Architecture)
 - REST endpoints can easily be generated
 - Each service has its own “global operation log” queue
 - services / global operation logs being eventually consistent

Response of GET /contacts/update-events?snapshot=<causal history>

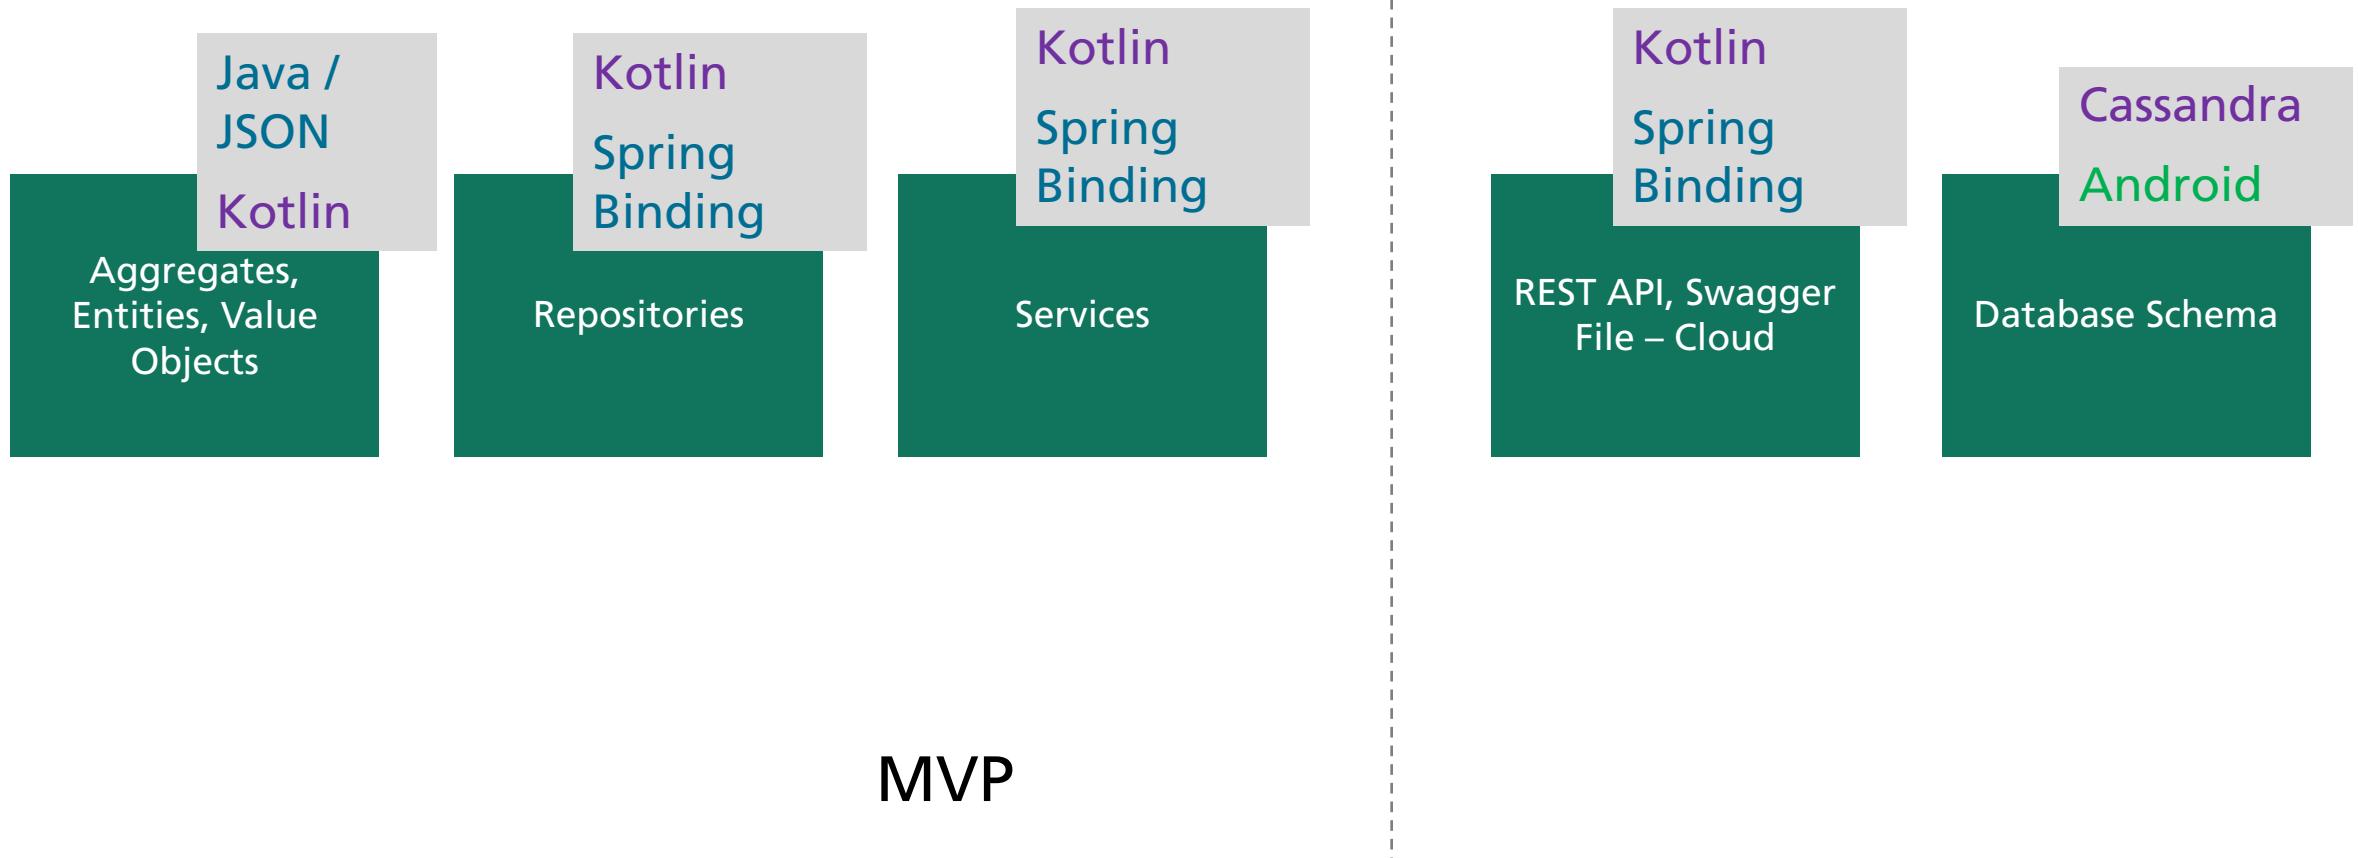
```
1  [
2  {
3      "id": "d10b642c-6e76-11e8-adc0-fa7ae01bbebc",
4      "readVersion": [ecfe3992-b0e4-11e8-96f8-529269fb1459, 080c2a0a-b0e5-11e8-96f8-529269fb1459, ...],
5
6      "operation": "de.fraunhofer.iese.oih.Contact.delete()", ←
7      "operationTime": "2007-04-05T12:30-02:00",
8      "operationOriginAppUuid": "com.snazzycontacts.SnazzyContacts",
9
10     "securityUserUuid": "bc9c46fe-238b-11e8-b467-0ed5f89f718b",
11     "securityUserRole": "some-role",
12
13     "parameters": []
14 },
15 {
16     "id": "7d2817e6-6e77-11e8-adc0-fa7ae01bbebc",
17     "readVersion": [ecfe3992-b0e4-11e8-96f8-529269fb1459, 080c2a0a-b0e5-11e8-96f8-529269fb1459, ...],
18
19     "operation": "de.fraunhofer.iese.oih.Contact.update(de.fraunhofer.iese.oih.Contact)", ←
20     "operationTime": "2007-04-05T12:30-02:00",
21     "operationOriginAppUuid": "com.snazzycontacts.SnazzyContacts",
22
23     "securityUserUuid": "bc9c46fe-238b-11e8-b467-0ed5f89f718b",
24     "securityUserRole": "some-role",
25
26     "parameters": [
27         {
28             "firstName": "Torsten",
29             "lastName": "Kuhn"
30         }
31     }
32 }
33 ]
34 ]
```

Causal Histories

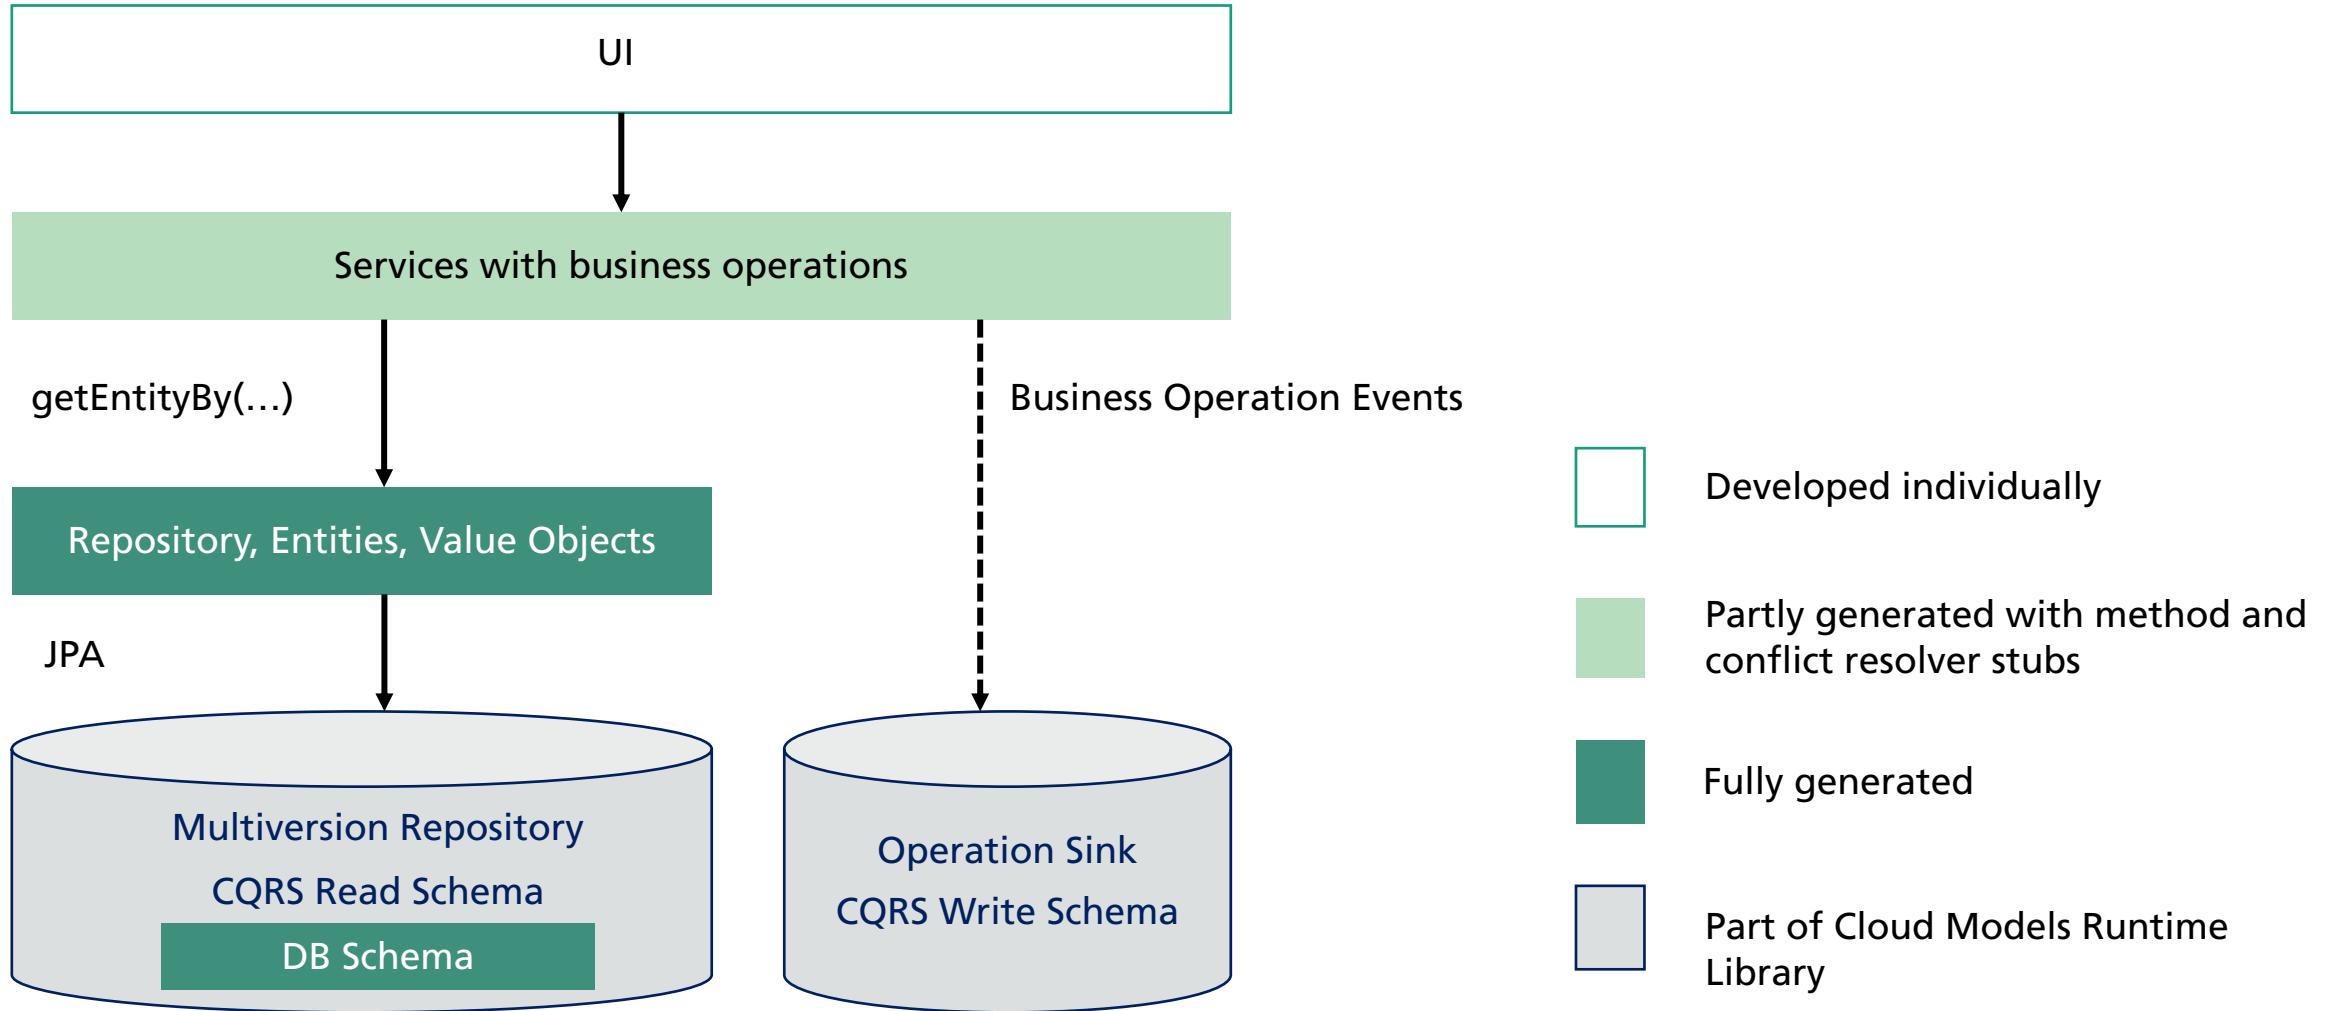
Operation Event Sink and Global Operation Log – Mobile Sync Scenario

- One (horizontally scalable) Central Service Replica in the Cloud functioning as a hub
 - designed similar to services of the cloud scenario from the previous slide
 - mobile apps use it to poll remote updates
 - mobile apps use it to transmit local updates
 - -> API of the central replica can again be generated (same approach as in cloud scenario)
- Apps have operation event sink and global operation log on the mobile device

Cloud Models Generated Replica Components

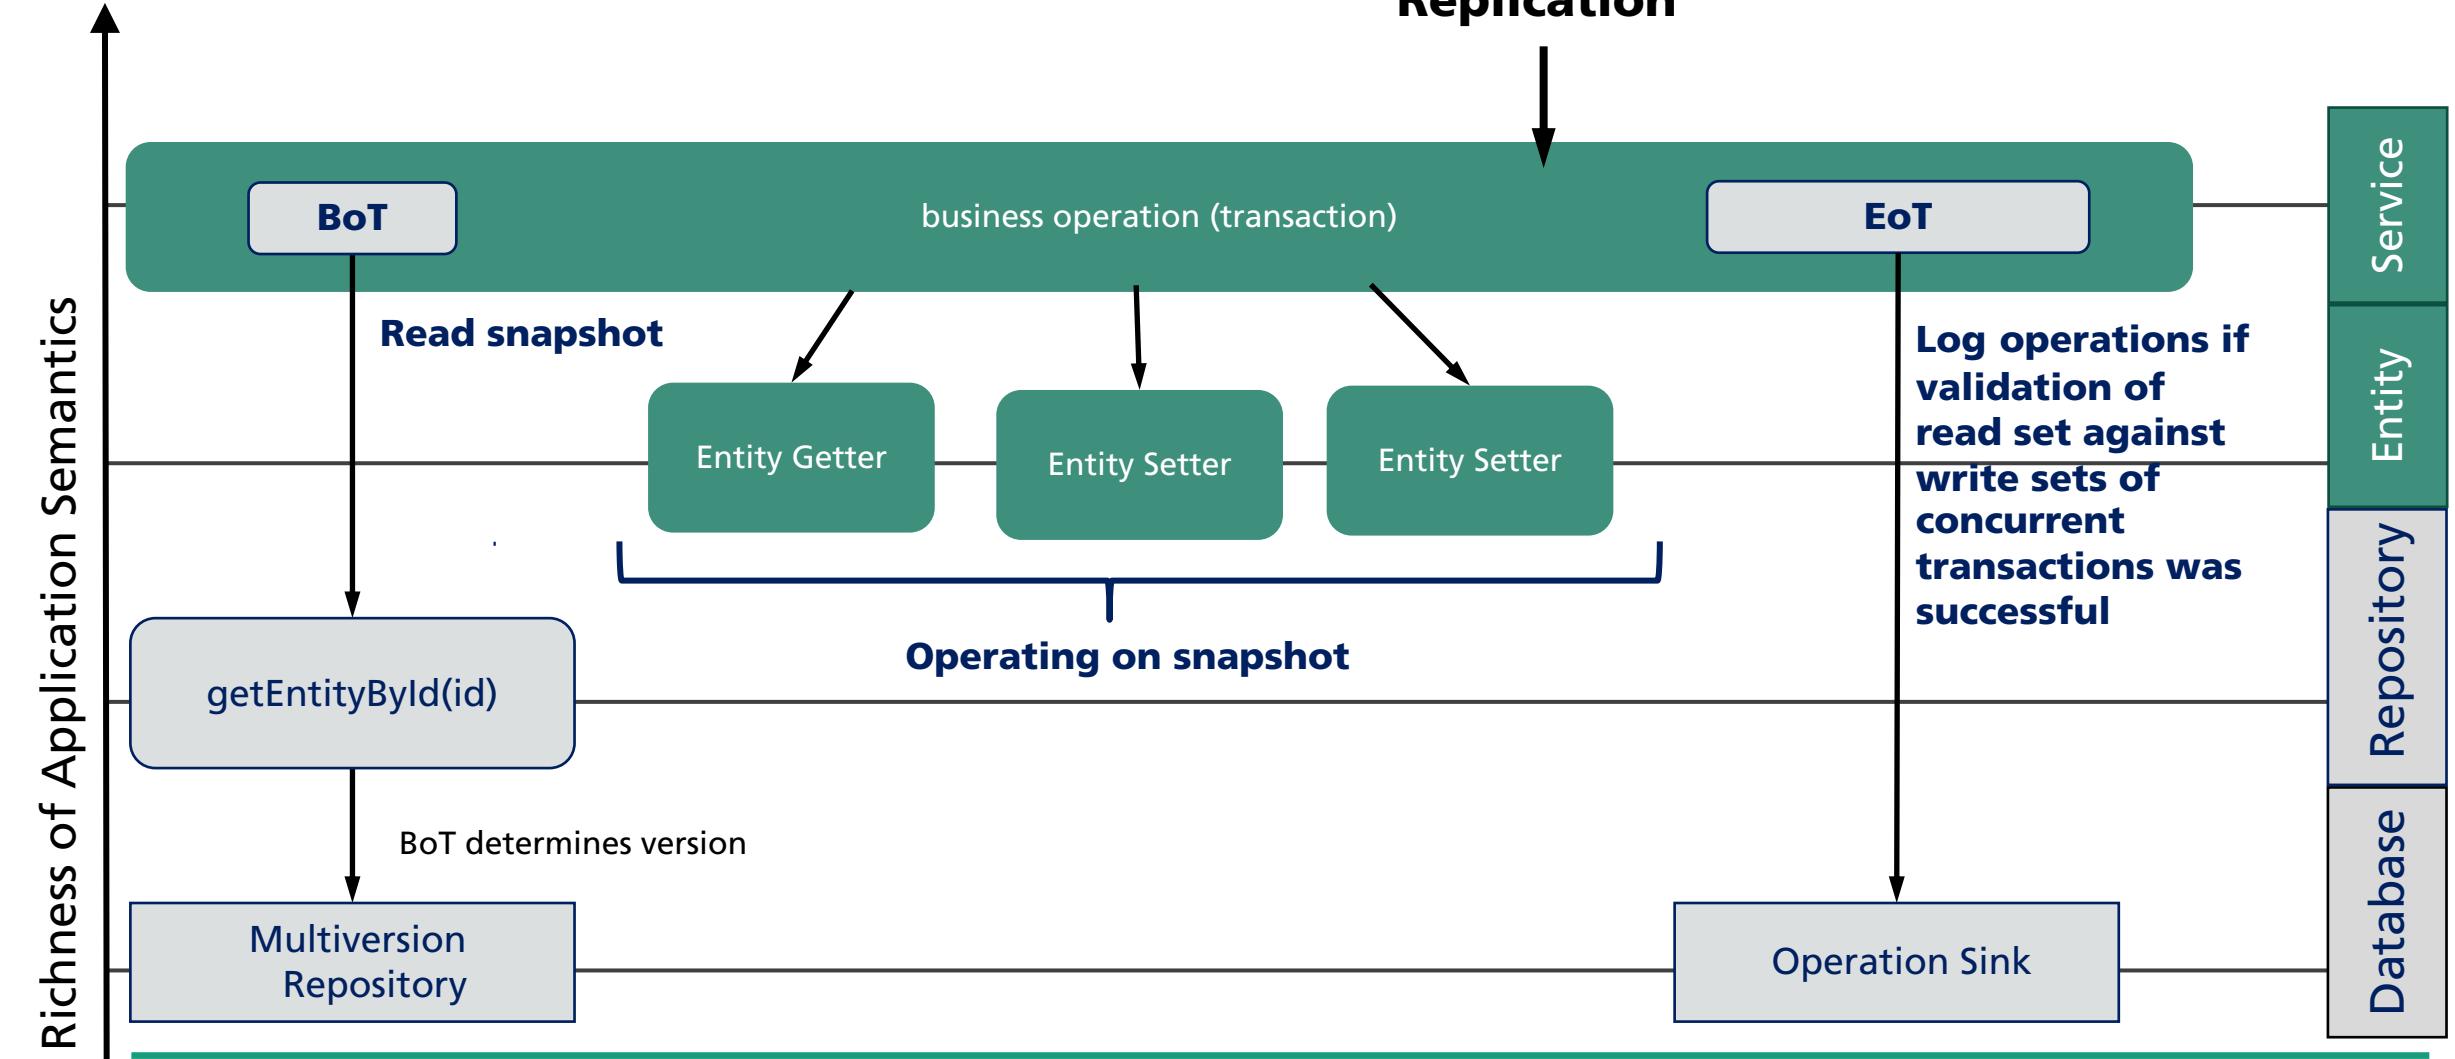


Layered Architecture of Microservices / Mobile Apps

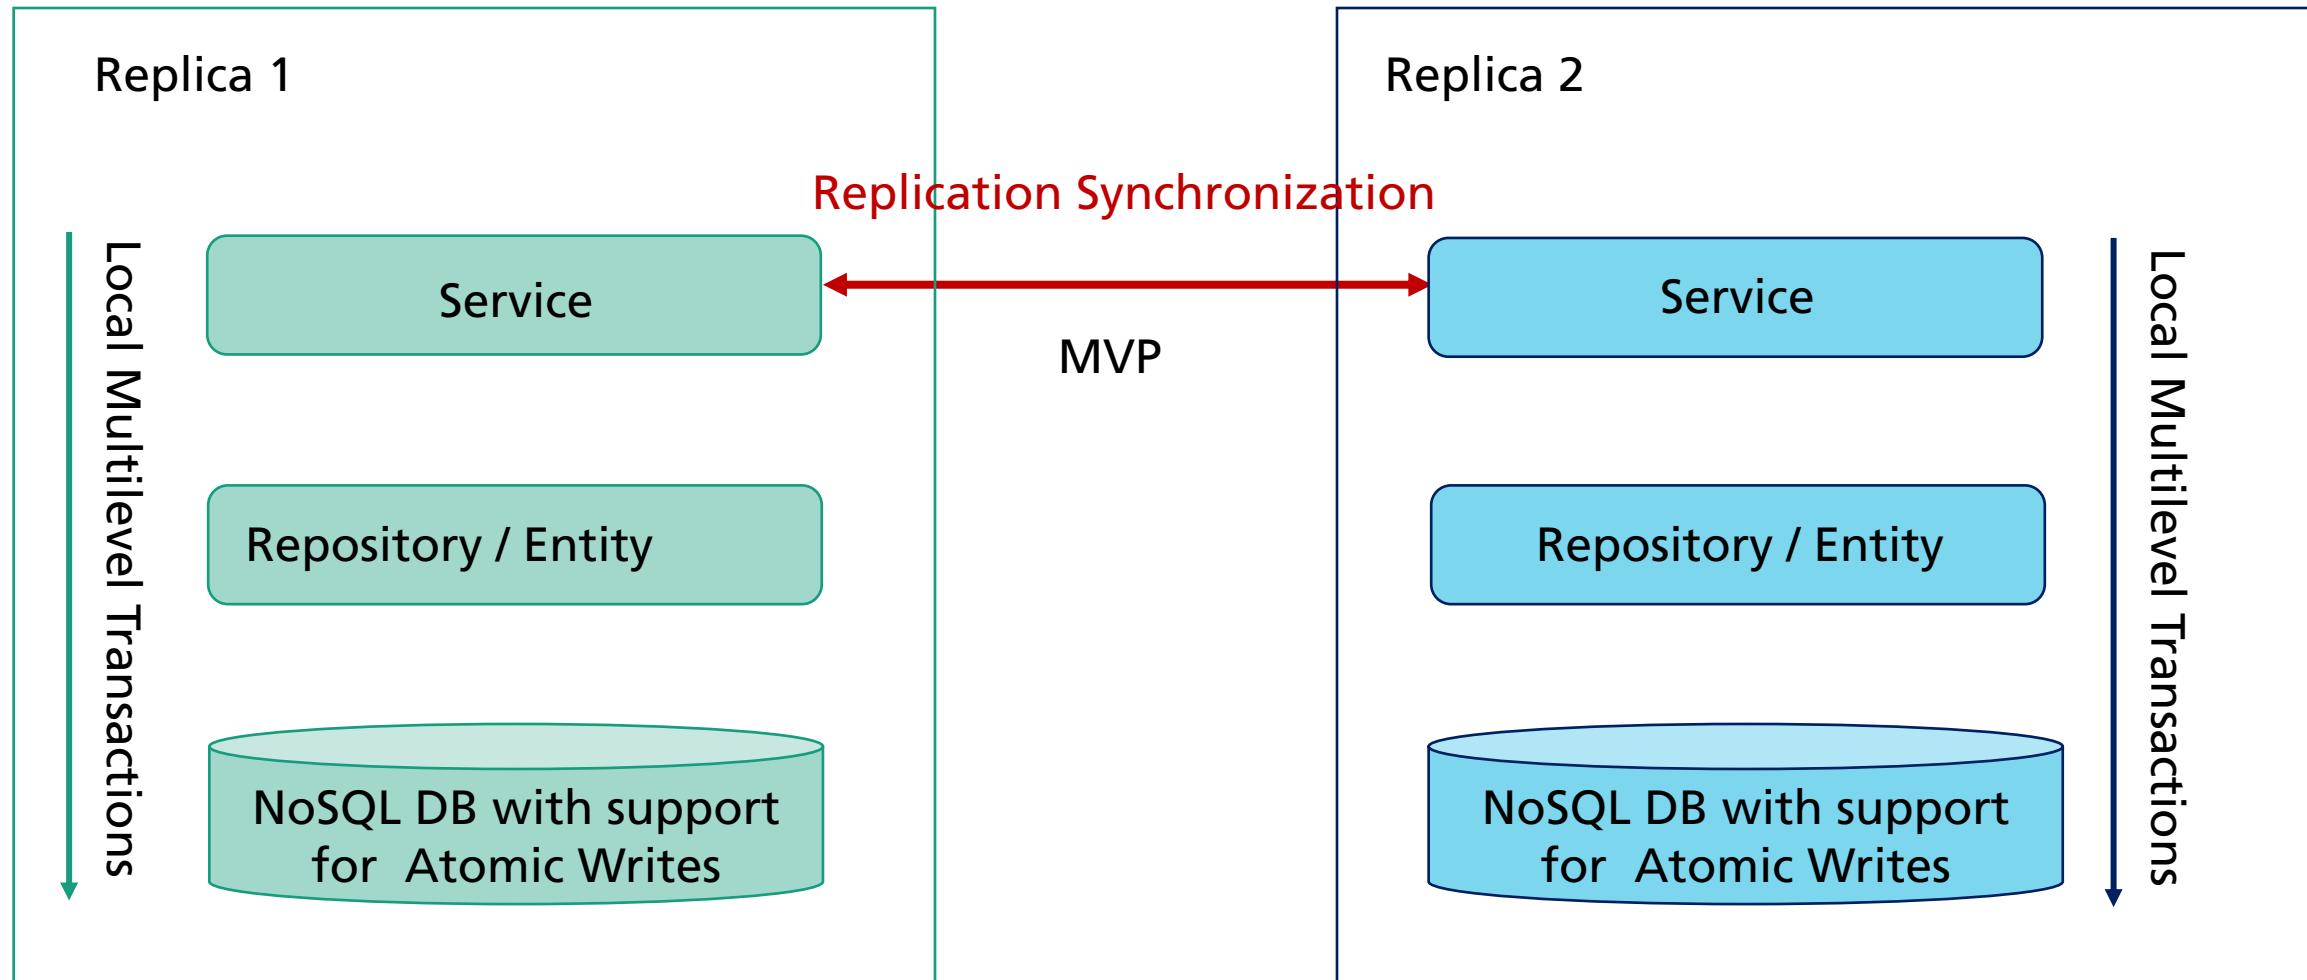


Level of Cloud Models Replication

Cloud Models Replication



Outlook: Cloud Models Multilevel Transactions



Cloud Models Testing Tools



Cloud Models Roadmap

- Ecosystem Data Sync: Lessons Learned, Best Practices, First Ideas & State of the Art Analysis -> John Deere Cooperation
- Cloud Models Core Library & Code Generator -> Software Campus Project (cloud-models.de), Digitale Teams Project (digitale-teams.de)
- Methodology & Design Guidelines -> PhD
- **OIH Connector for Cloud Model Apps -> OIH Project**

Integration Issues & Challenges OIH Connector

- OIH /data REST endpoint can be used and data can be taken directly from the multi-version repository of the cloud-models framework
 - How is versioning treated in OIH ?
 - -> Connector will maintain pending operations that need to be scheduled for synchronization with OIH
- Remote updates polled from /webhooks endpoint
 - Problem: No operations can be queried via /webhooks endpoint
 - Again: How is versioning treated in OIH ?
 - CRUD operations need to be derived by the connector
 - **how to determine the identity of a (derived) CRUD operation ?**
 - **what happens if a update-events are delivered more than once?**
 - **how to prevent "update cycles"?**
 - remote updates events need to be treated differently, than locale ones that are scheduled for propagation
 - Cloud Models Replication can only run on level of basic CRUD operations. No support for high-level service operations. No sophisticated conflict handling.
 - **Adaption of the cloud models framework necessary**, s.t. replication level is configurable (**basic CRUD operations vs. service operations (preferred)**)
 - **Conflicts will be resolved by the replicas with the deterministic conflict resolution algorithm of cloud models framework (based on the conflict relationships of basic read-and-write operations)**
 - **cloud model apps will be able to register conflict resolvers via DSL of the cloud models framework**
 - **Consequence: cloud models apps might not be eventually consistent with OIH apps, in case OIH applies a different reconciliation algorithm.**



Susanne Braun

Tel.: +49 631 6800 2138

susanne.braun@iese.fraunhofer.de

Twitter: @susannebraun